



# Network Apps: Overview of Socket API

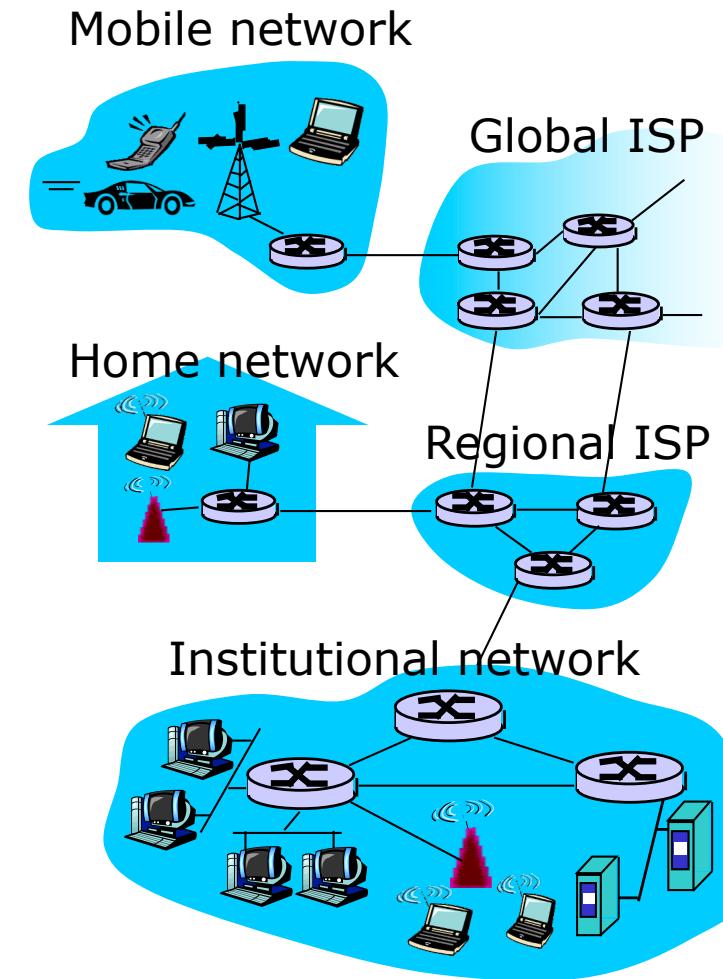
# Network Apps: HTTP & Content Delivery

Michael Kirkedal Thomsen

Based on slides compiled by Marcos Vaz Salles; With adaptations by  
Vivek Shah

# Recap: What was the main topics

- What was the main topics last week?



Source: Kurose & Ross



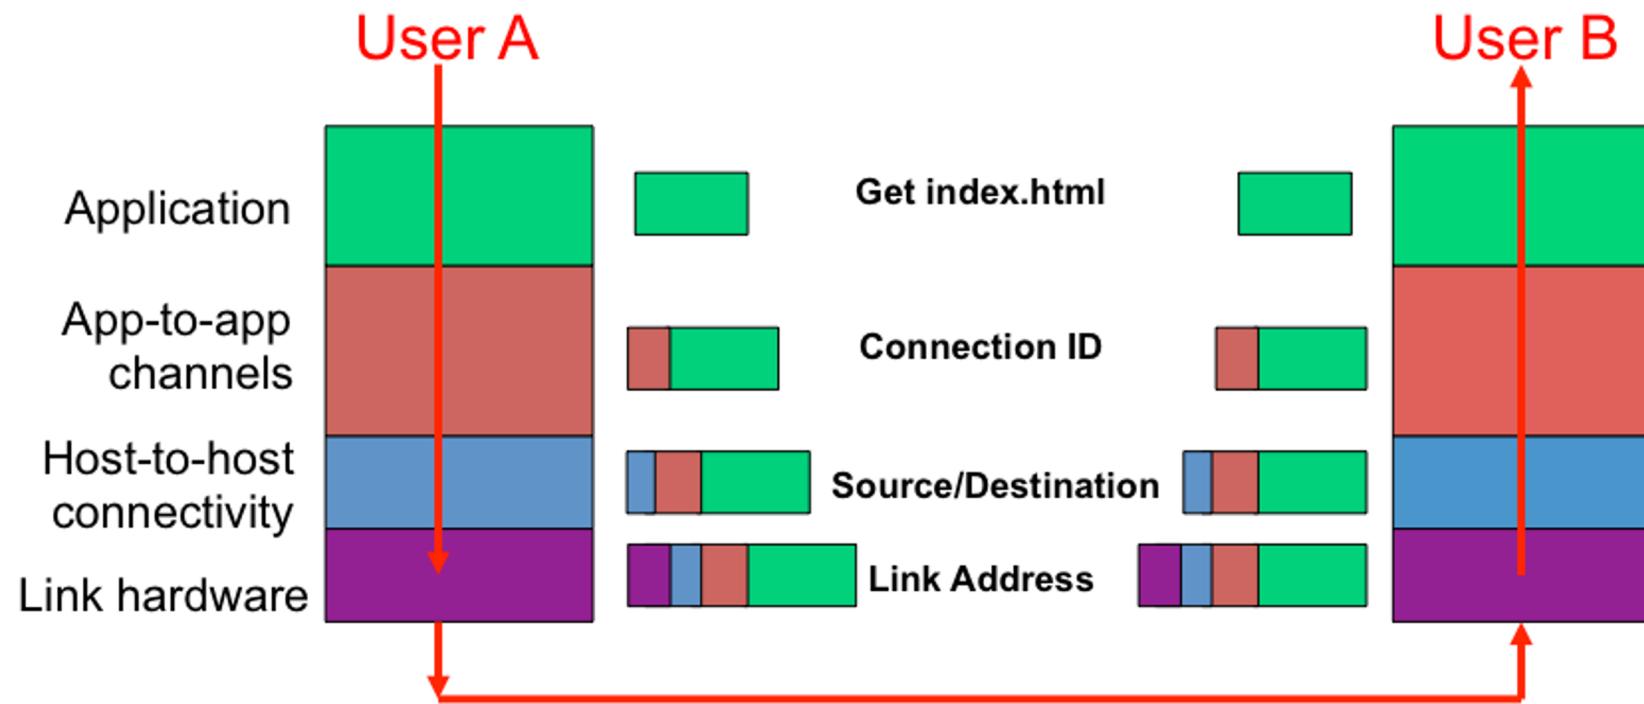
## Recap: Key Concepts in Networking

- **Protocols**
  - Speaking the same language
  - Syntax and semantics
- **Layering**
  - Standing on the shoulders of giants
  - A key to managing complexity
- **Resource allocation**
  - Dividing scarce resources among competing parties
  - Memory, link bandwidth, wireless spectrum, paths
- **Naming**
  - What to call computers, services, protocols, ...

Source: Freedman



## Recap: 5-Layer Encapsulation (OSI model)

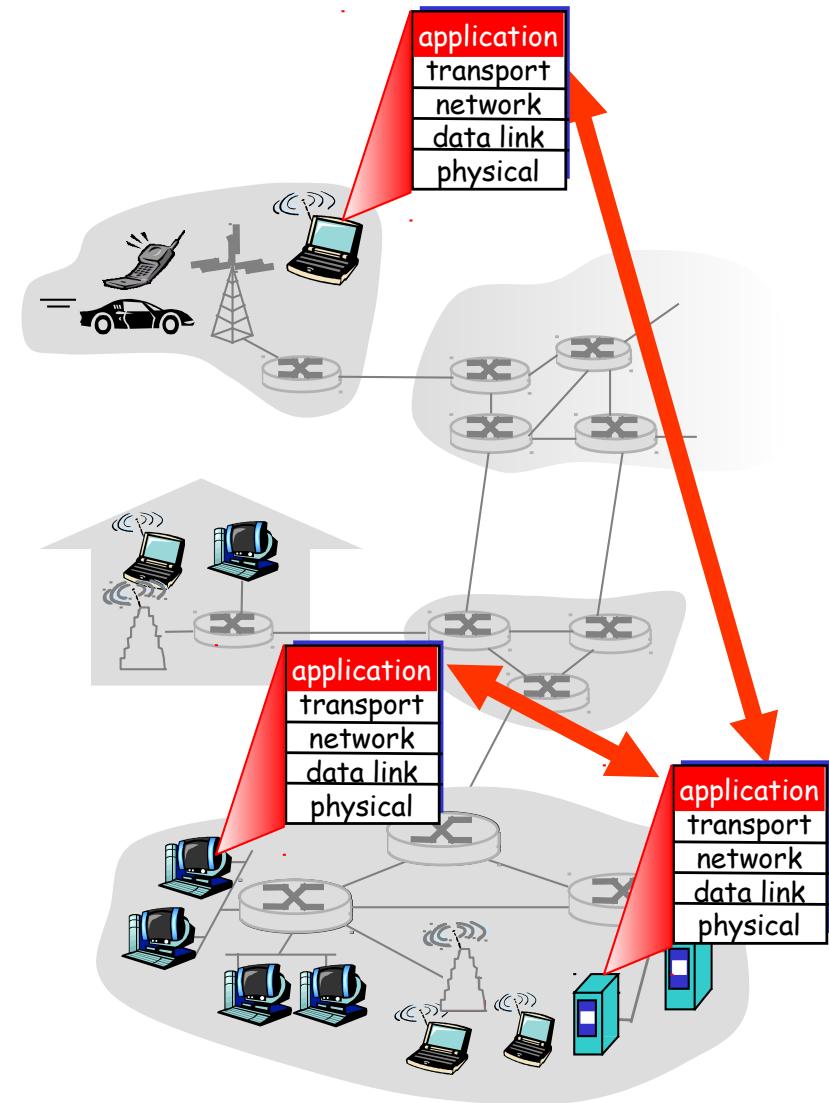


Source: Freedman



## Creating a network application

- write programs that
  - run on (different) end systems
  - communicate over network
  - e.g., web server software communicates with browser software
- No need to write software for network-core devices
  - network-core devices do not run user applications
  - applications on end systems allows for rapid app development, propagation



Source: Freedman



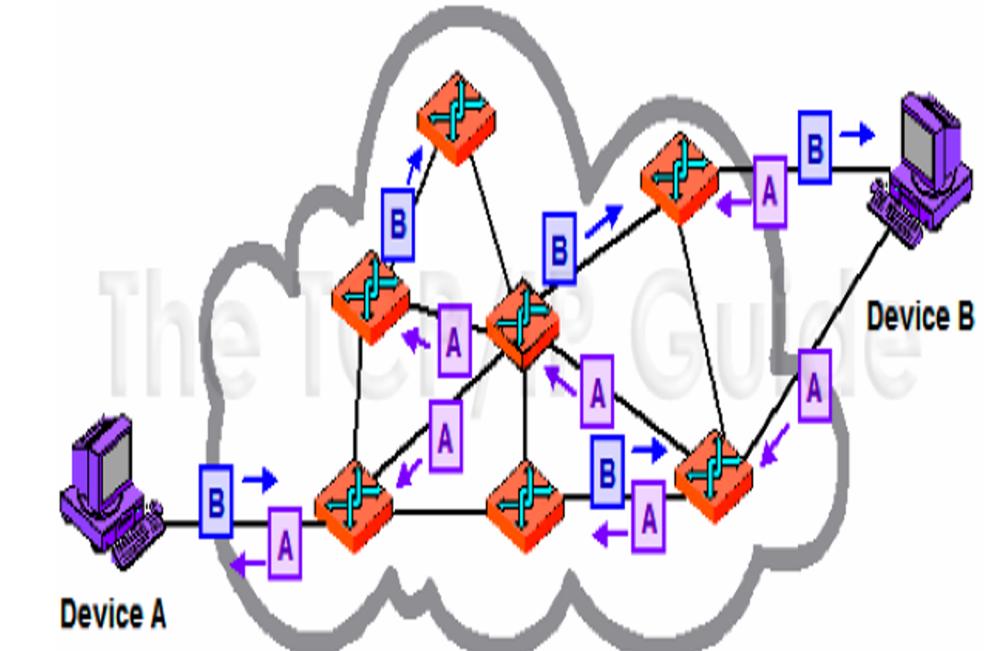
## Some network apps

- e-mail
- remote login
- web
- instant messaging
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube)
- voice over IP
- real-time video conferencing
- social networking
- cloud computing
- ...
- ...
- 



Network applications need streams of data

Circuit switching → Packet switching



[http://www.tcpipguide.com/free/t\\_CircuitSwitchingandPacketSwitchingNetworks-2.htm](http://www.tcpipguide.com/free/t_CircuitSwitchingandPacketSwitchingNetworks-2.htm)

Today's networks provide packet delivery, not streams!

Source: Freedman (partial)



## What if the Data Doesn't Fit?

GET /courses/archive/spr09/cos461/ HTTP/1.1

Host: www.cs.princeton.edu

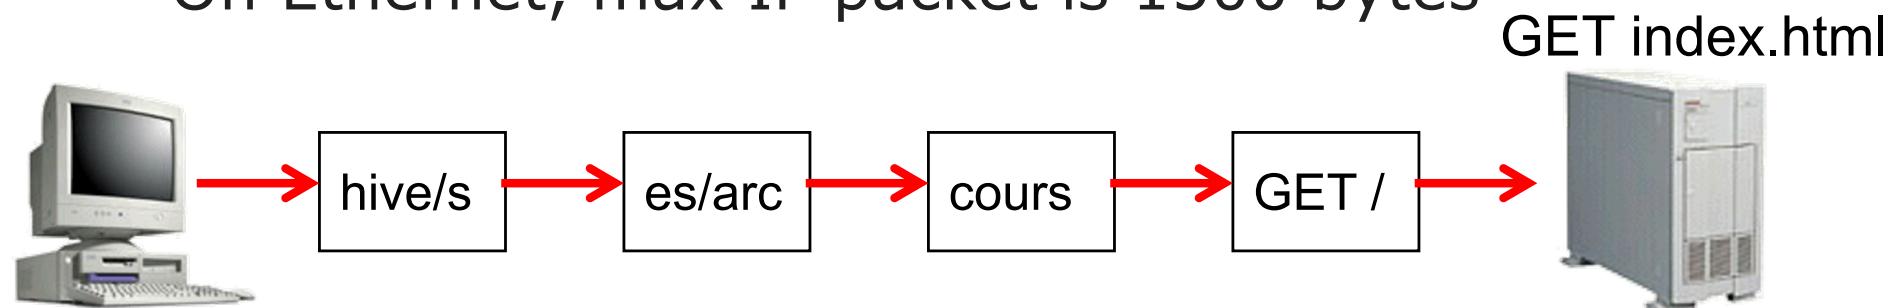
User-Agent: Mozilla/4.03

CRLF

### Request

#### Problem: Packet size

- Typical Web page is 10 kbytes
- On Ethernet, max IP packet is 1500 bytes



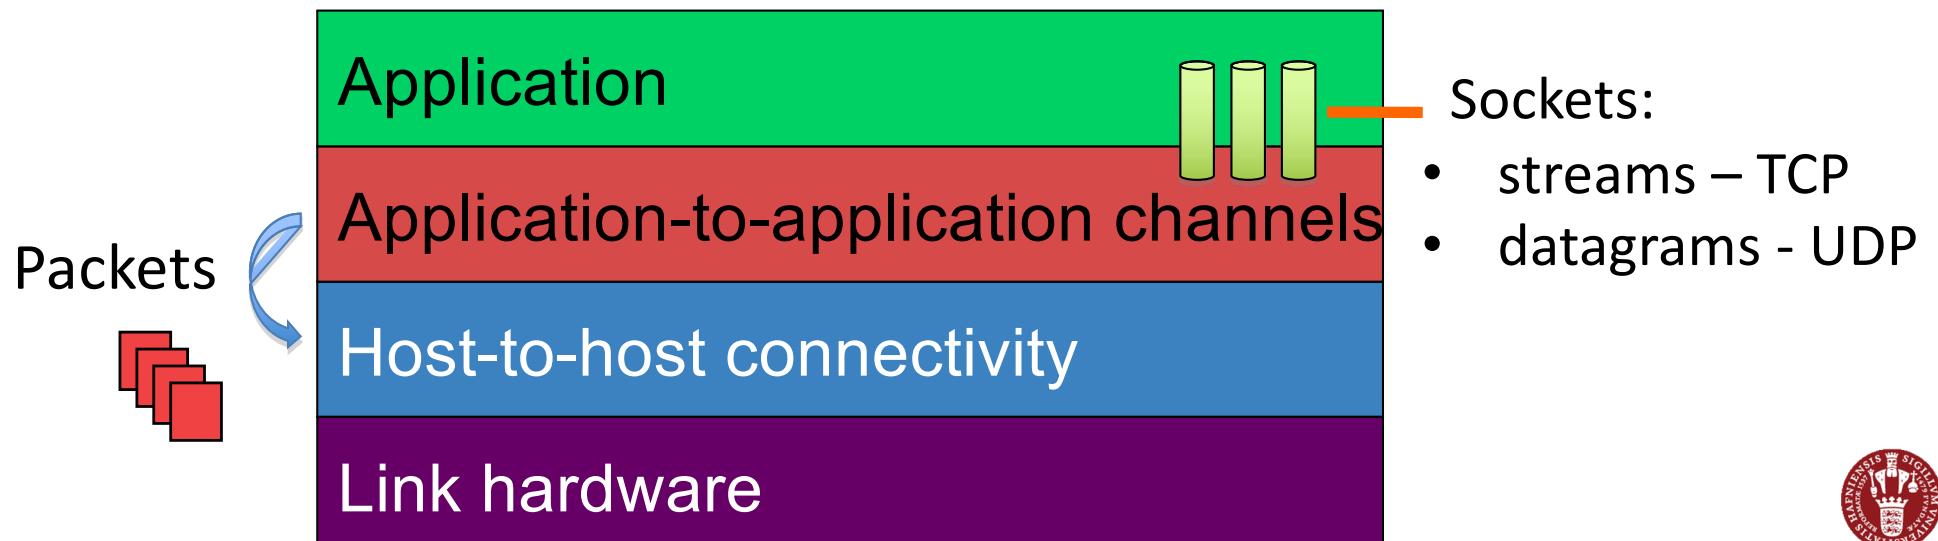
#### Solution: Split the data across multiple packets

Source: Freedman



# Layering = Functional Abstraction

- Sub-divide the problem
  - Each layer relies on services from layer below
  - Each layer exports services to layer above
- Interface between layers defines interaction
  - Hides implementation details
  - Layers can change without disturbing other layers

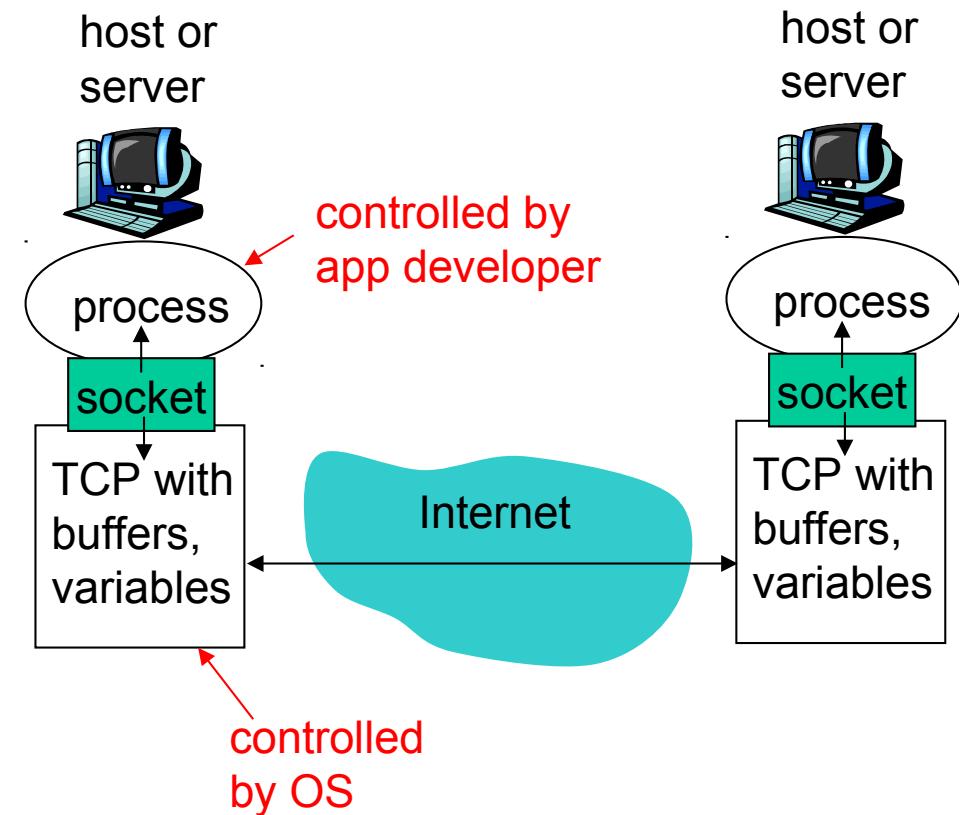


Source: Freedman



## Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
  - sending process shoves message out door
  - sending process relies on transport infrastructure on other side of door which brings message to socket at receiving process



API: (1) choice of transport protocol; (2) ability to fix a few parameters (more on this in next lecture!)

Source: Freedman



# Internet transport protocols services

## TCP service:

- *connection-oriented*: setup required between client and server processes
- *reliable transport*: between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantees, security

## UDP service:

- unreliable data transfer between sending and receiving process
- does not provide: connection setup, reliability, flow control, congestion control, timing, throughput guarantee, or security

Q: why bother? Why is there a UDP?



Source: Kurose & Ross

# UNIX Socket API

- Socket interface
  - Originally provided in Berkeley UNIX
  - Later adopted by all popular operating systems
  - Simplifies porting applications to different OSes
- In UNIX, everything is like a file
  - All input is like reading a file
  - All output is like writing a file
  - File is represented by an integer file descriptor
- API implemented as system calls
  - E.g., connect, read, write, close, ...

Source: Freedman

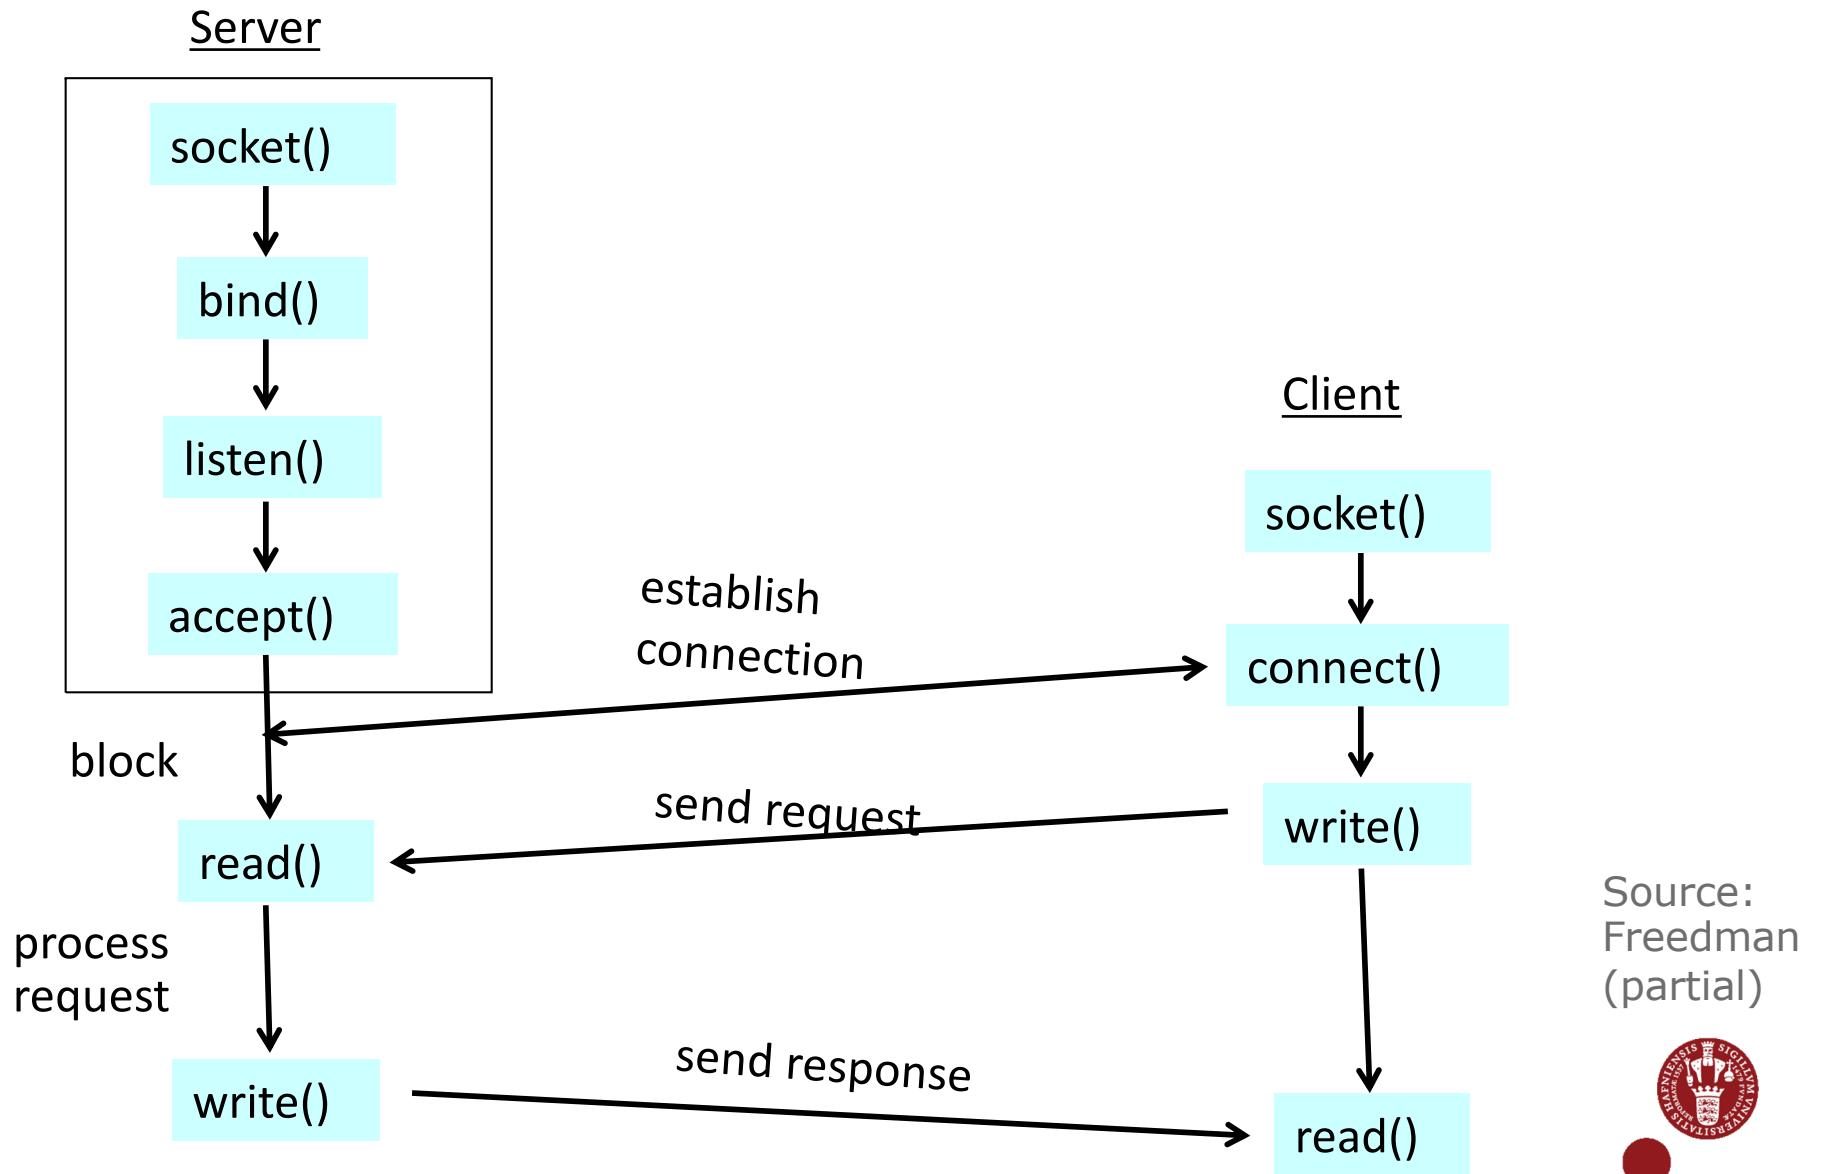


# Identifying the Receiving Process

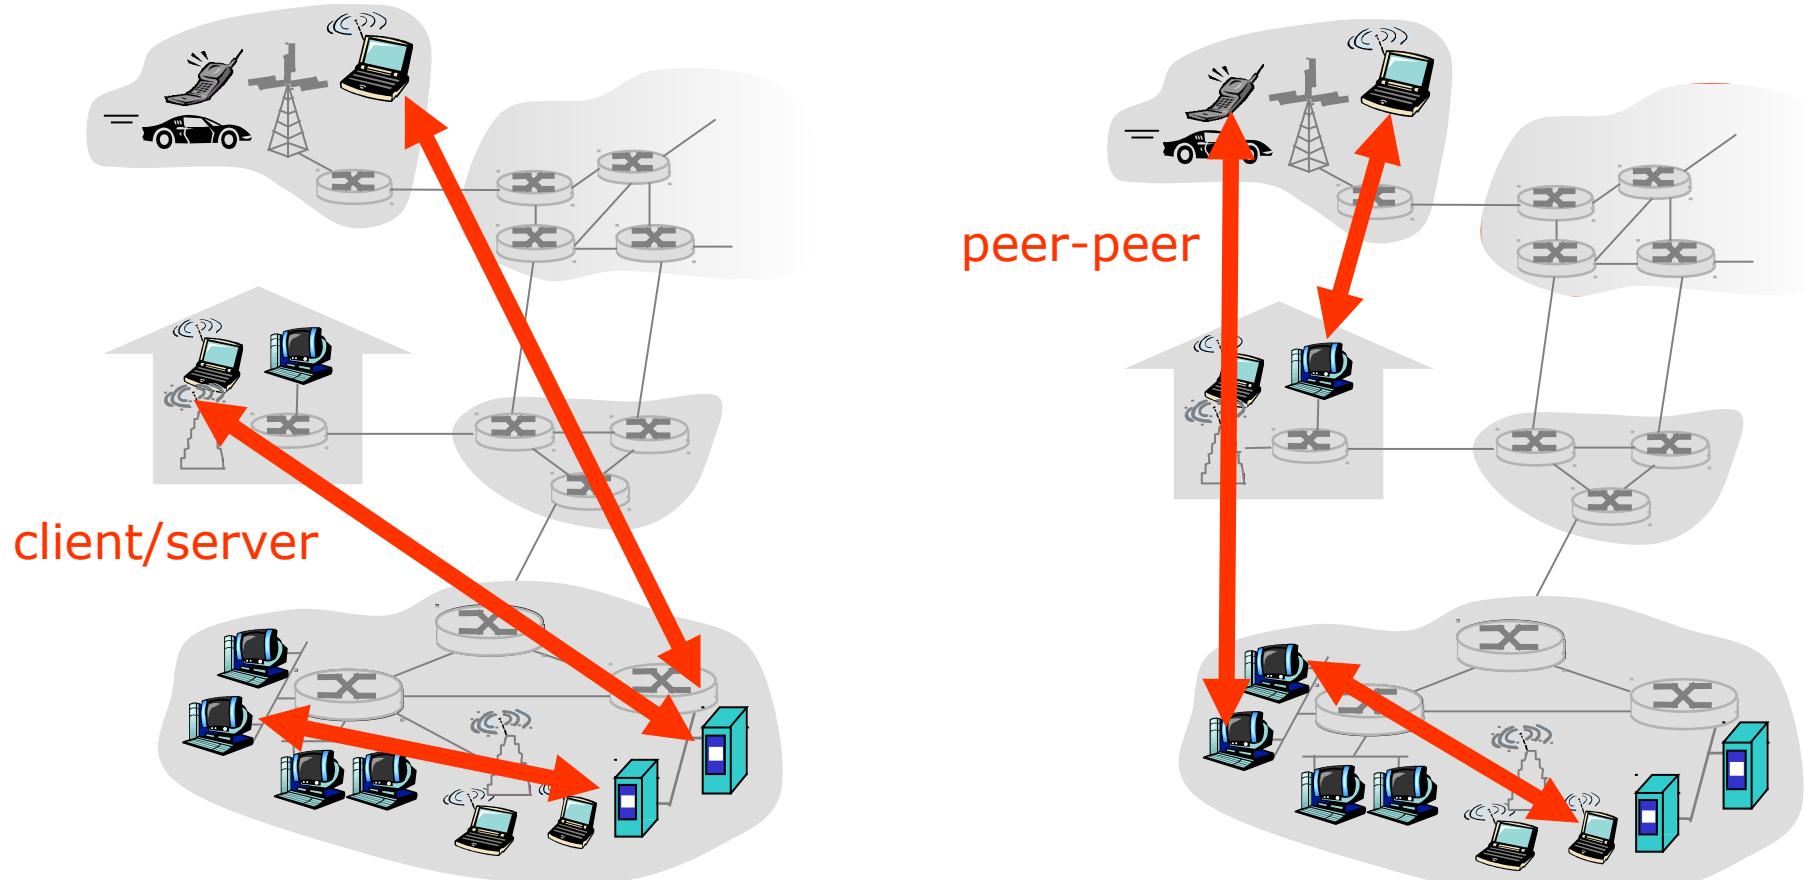
- Sending process must identify the receiver
  - The receiving end host machine
  - The specific socket in a process on that machine
- Receiving host
  - Destination address that uniquely identifies the host
  - Typically, high-level name translated to IP address (DNS)
  - For example, [www.diku.dk](http://www.diku.dk) → 130.225.96.108
  - An IP address is a 32-bit quantity
- Receiving socket
  - Host may be running many different processes
  - Destination port that uniquely identifies the socket
  - A port number is a 16-bit quantity



# Client-Server TCP Sockets



# Application Architectures

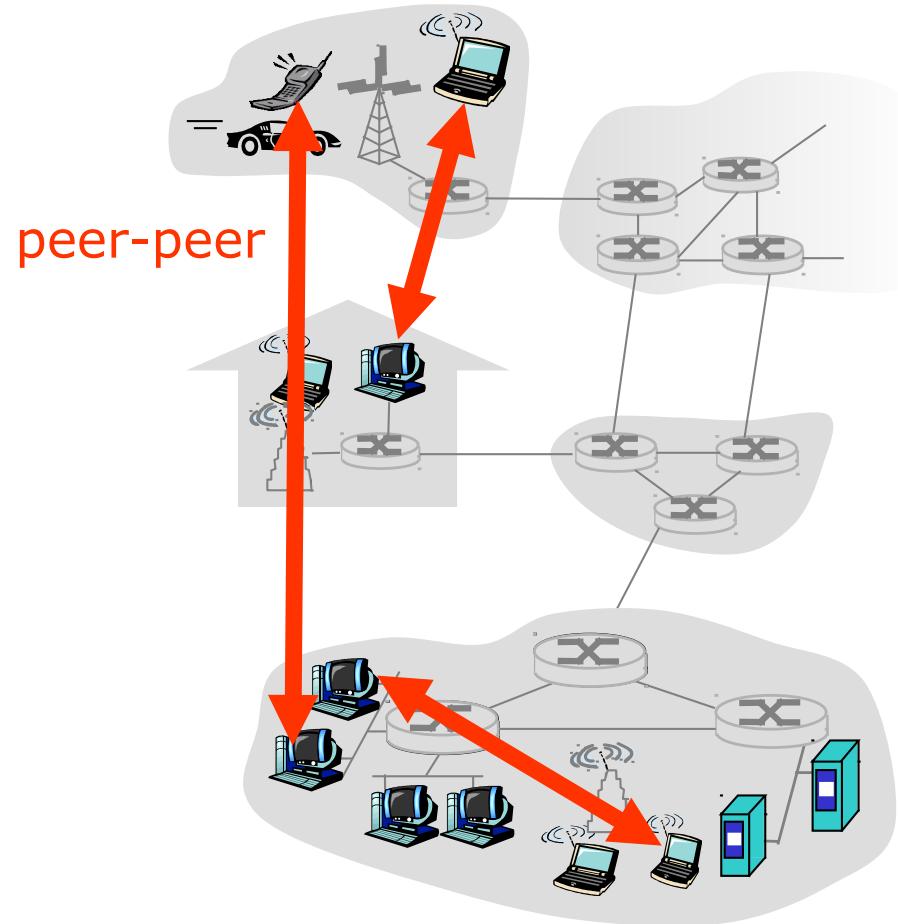
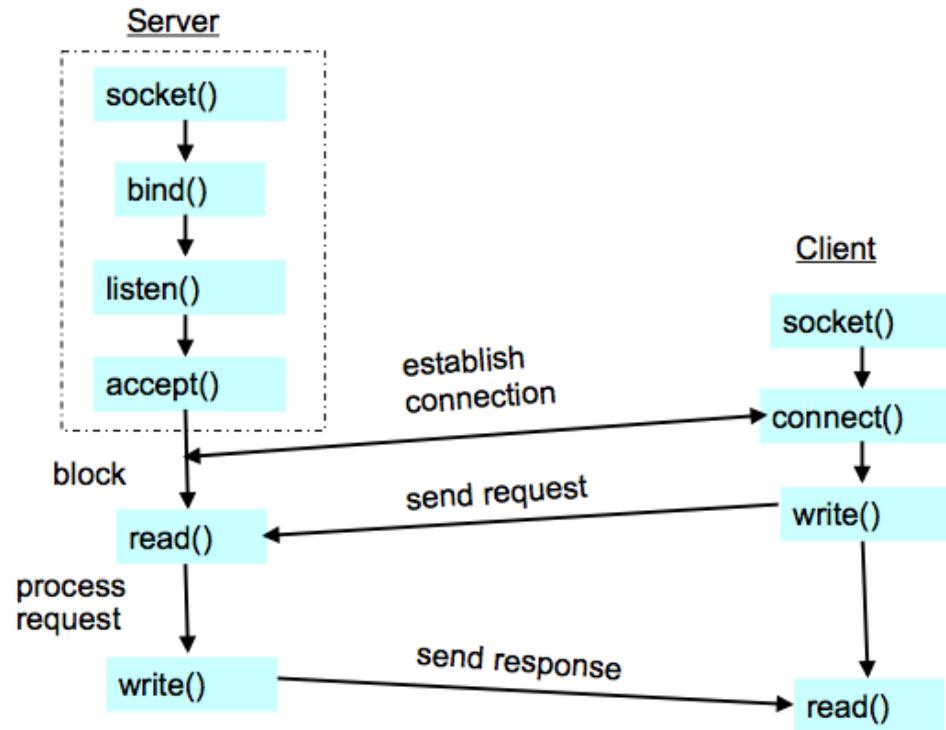


- P2P highly scalable, but difficult to manage
- Hybrids also possible, e.g., Skype

Source: Freedman (partial)



# Discussion: How do you set sockets up in a P2P application?



- How many sockets in each peer?
  - What if many peers connect to one peer?

Source: Freedman (partial)



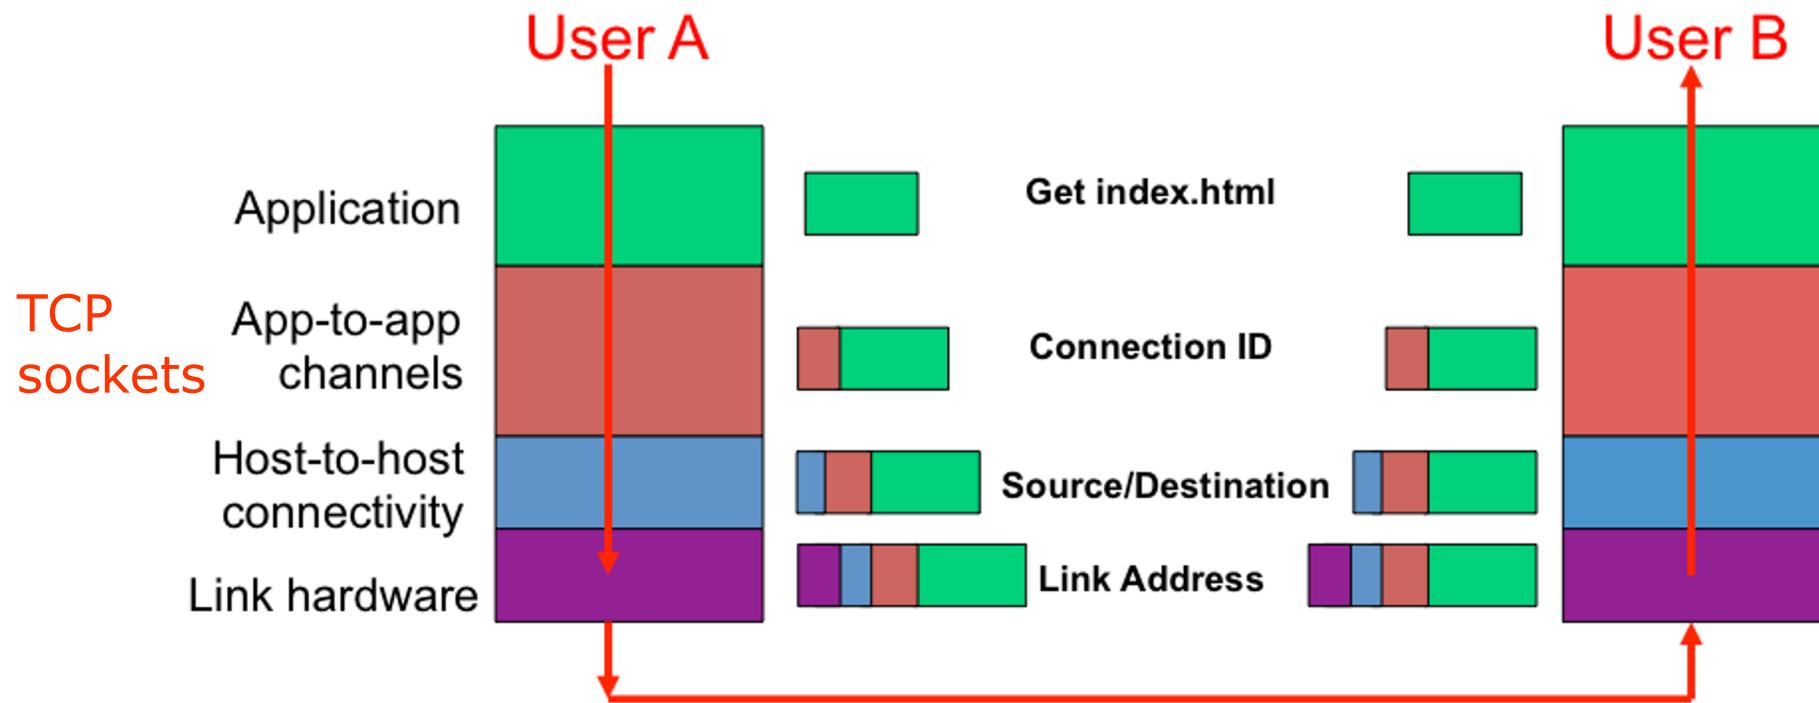
# HTTP Basics

- HTTP layered over bidirectional byte stream
- Interaction
  - Client sends request to server, followed by response from server to client
  - Requests/responses are encoded in text
- Targets access to web objects
  - GET, POST, HEAD → HTTP/1.0
  - GET, POST, HEAD, PUT, DELETE → HTTP/1.1
- Stateless
  - Server maintains no info about past client requests
  - What about personalization? Data stored in back-end database; client sends “web cookie” used to lookup data



Source: Freedman (partial)

## Layer Encapsulation in HTTP



Source: Freedman (partial)



# HTTP Request Example

GET / HTTP/1.1

Host: sns.cs.princeton.edu

Accept: \*/\*

Accept-Language: en-us

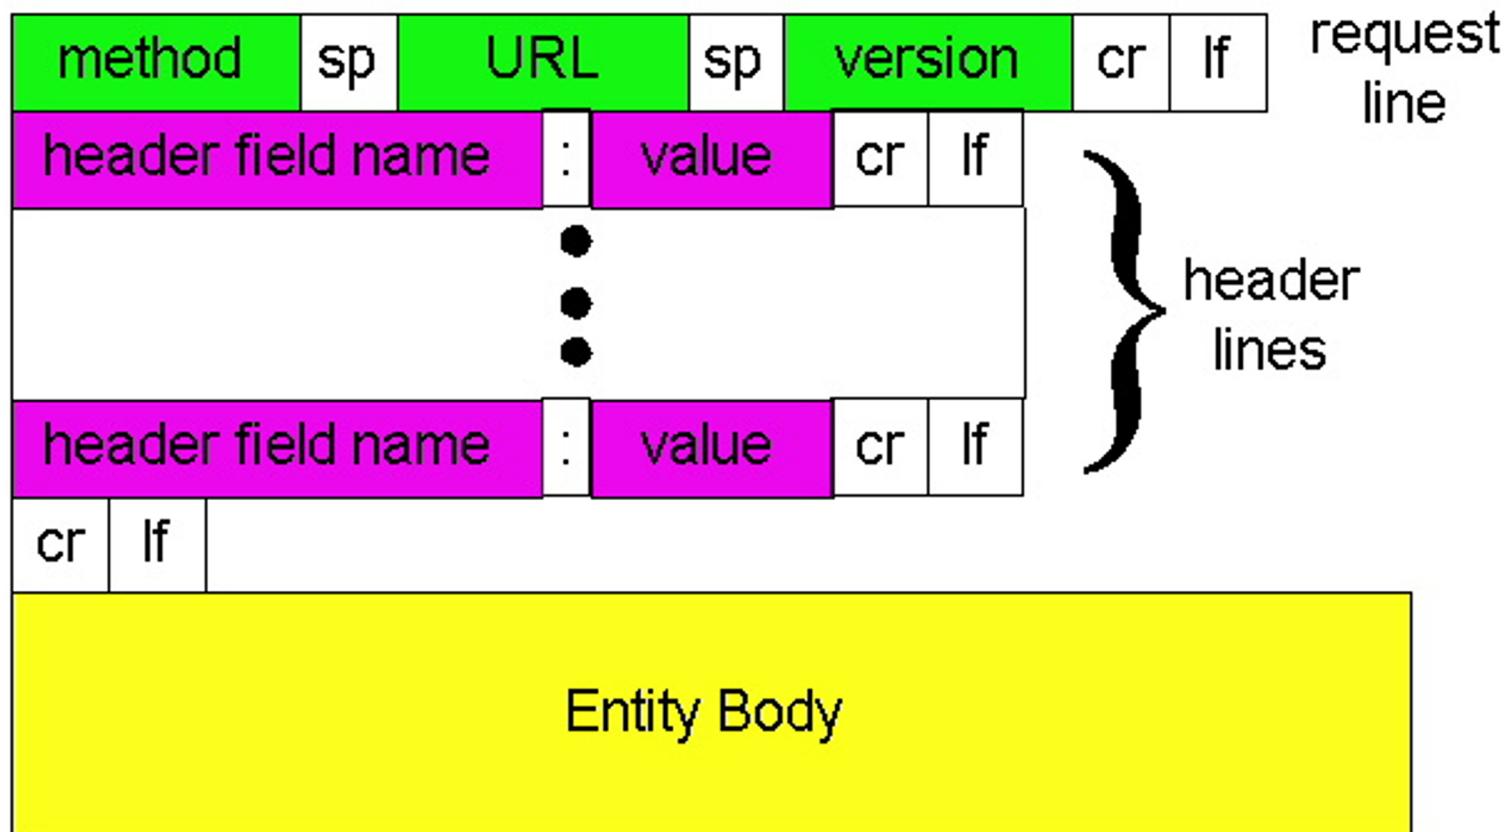
Accept-Encoding: gzip, deflate

User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X 10.5; en-US;  
rv:1.9.2.13) Gecko/20101203 Firefox/3.6.13

Connection: Keep-Alive



# HTTP Request



Source: Freedman



# HTTP Response Example

HTTP/1.1 200 OK

Date: Wed, 02 Feb 2011 04:01:21 GMT

Server: Apache/2.2.3 (CentOS)

X-Pingback: http://sns.cs.princeton.edu/xmlrpc.php

Last-Modified: Wed, 01 Feb 2011 12:41:51 GMT

ETag: "7a11f-10ed-3a75ae4a"

Accept-Ranges: bytes

Content-Length: 4333

Keep-Alive: timeout=15, max=100

Connection: Keep-Alive

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN"  
    "http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">  
<html xmlns="http://www.w3.org/1999/xhtml" dir="ltr" lang="en-US">
```

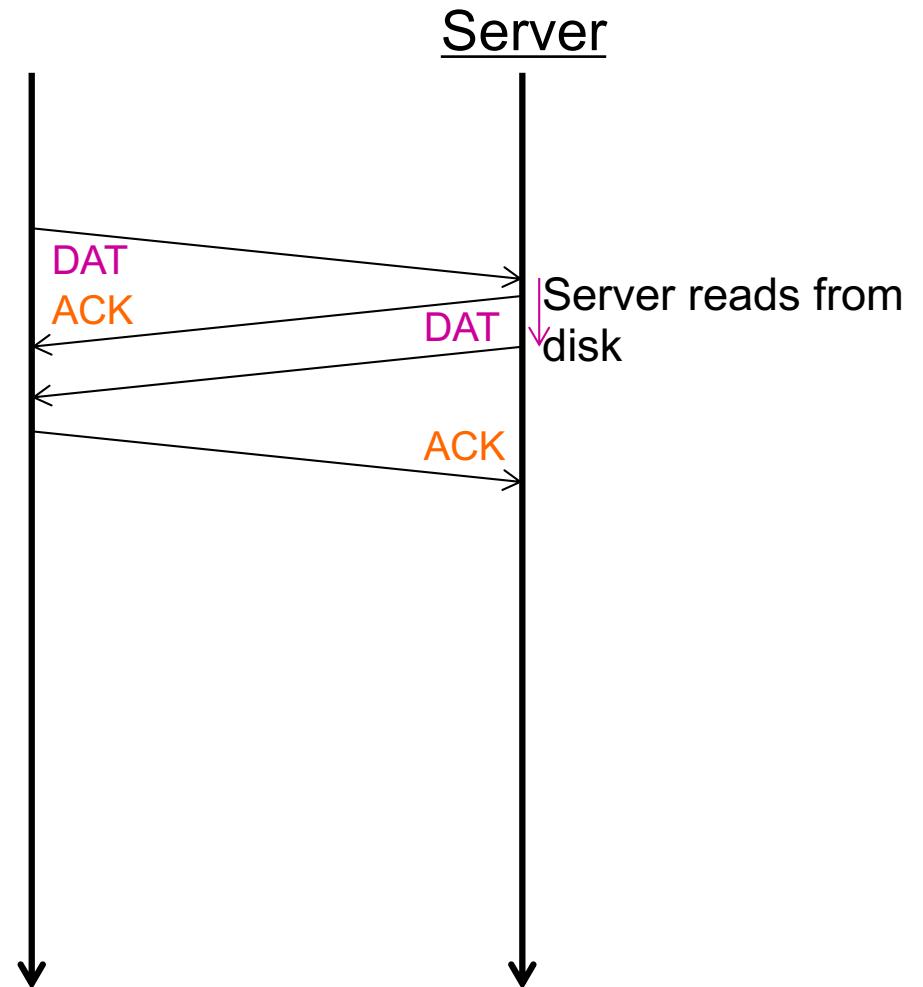


# Single Transfer Example

Source: Freedman

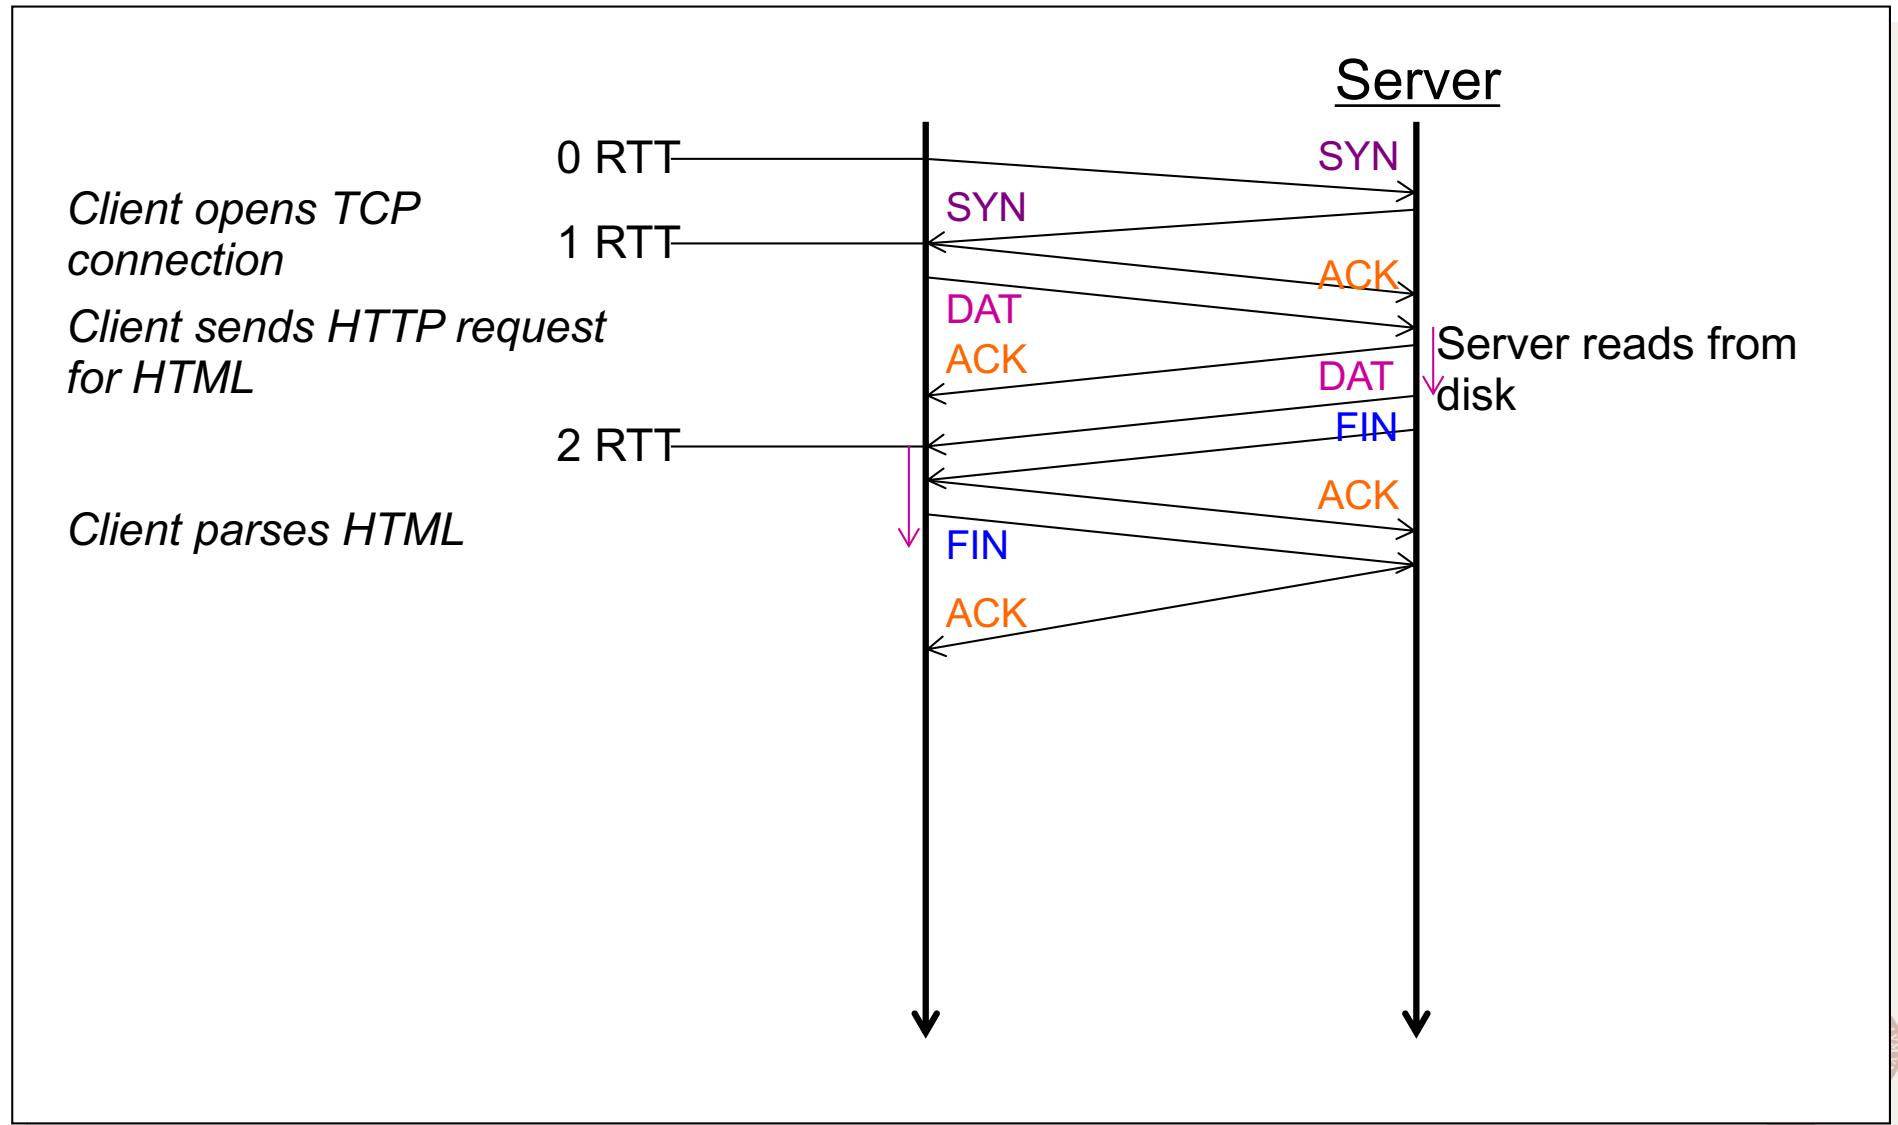
*Client sends HTTP request  
for HTML*

*Client parses HTML*



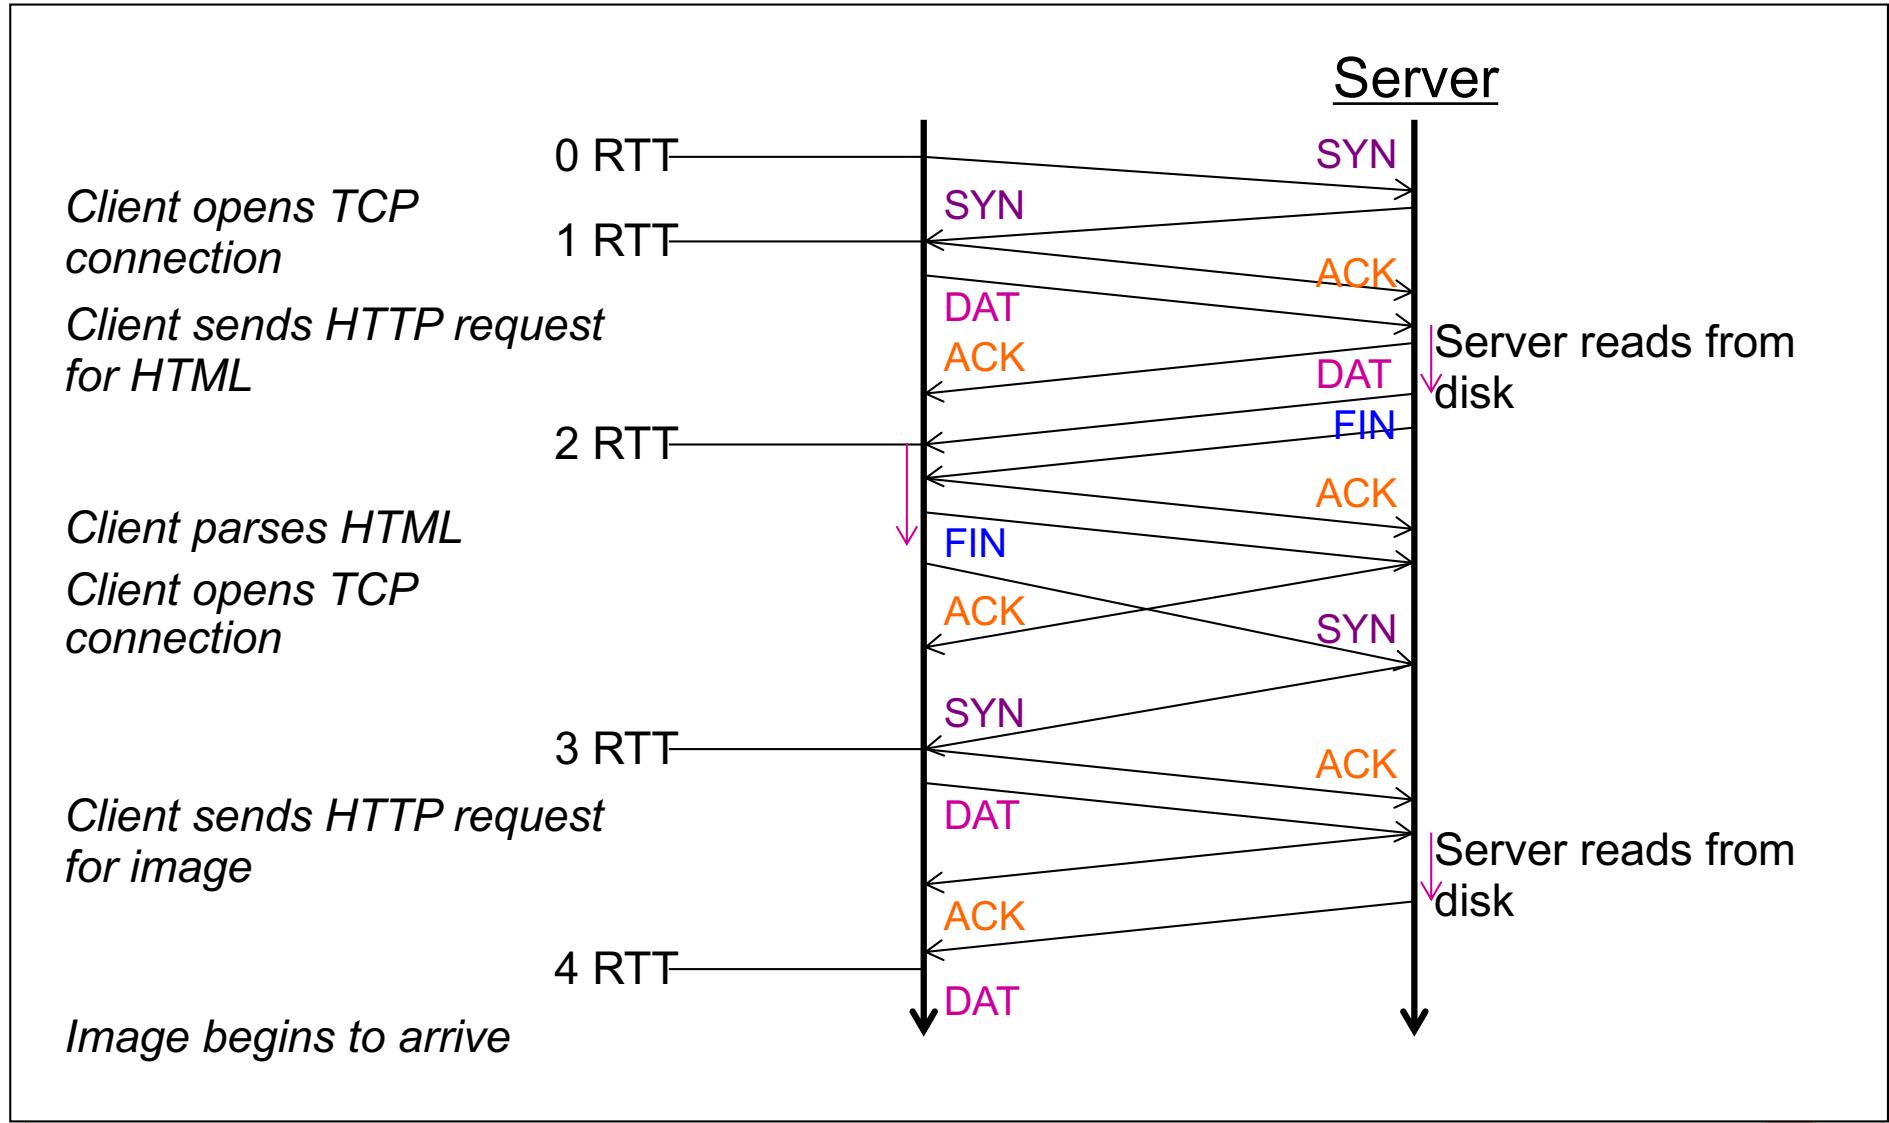
# Single Transfer Example

Source: Freedman



# Single Transfer Example

Source: Freedman



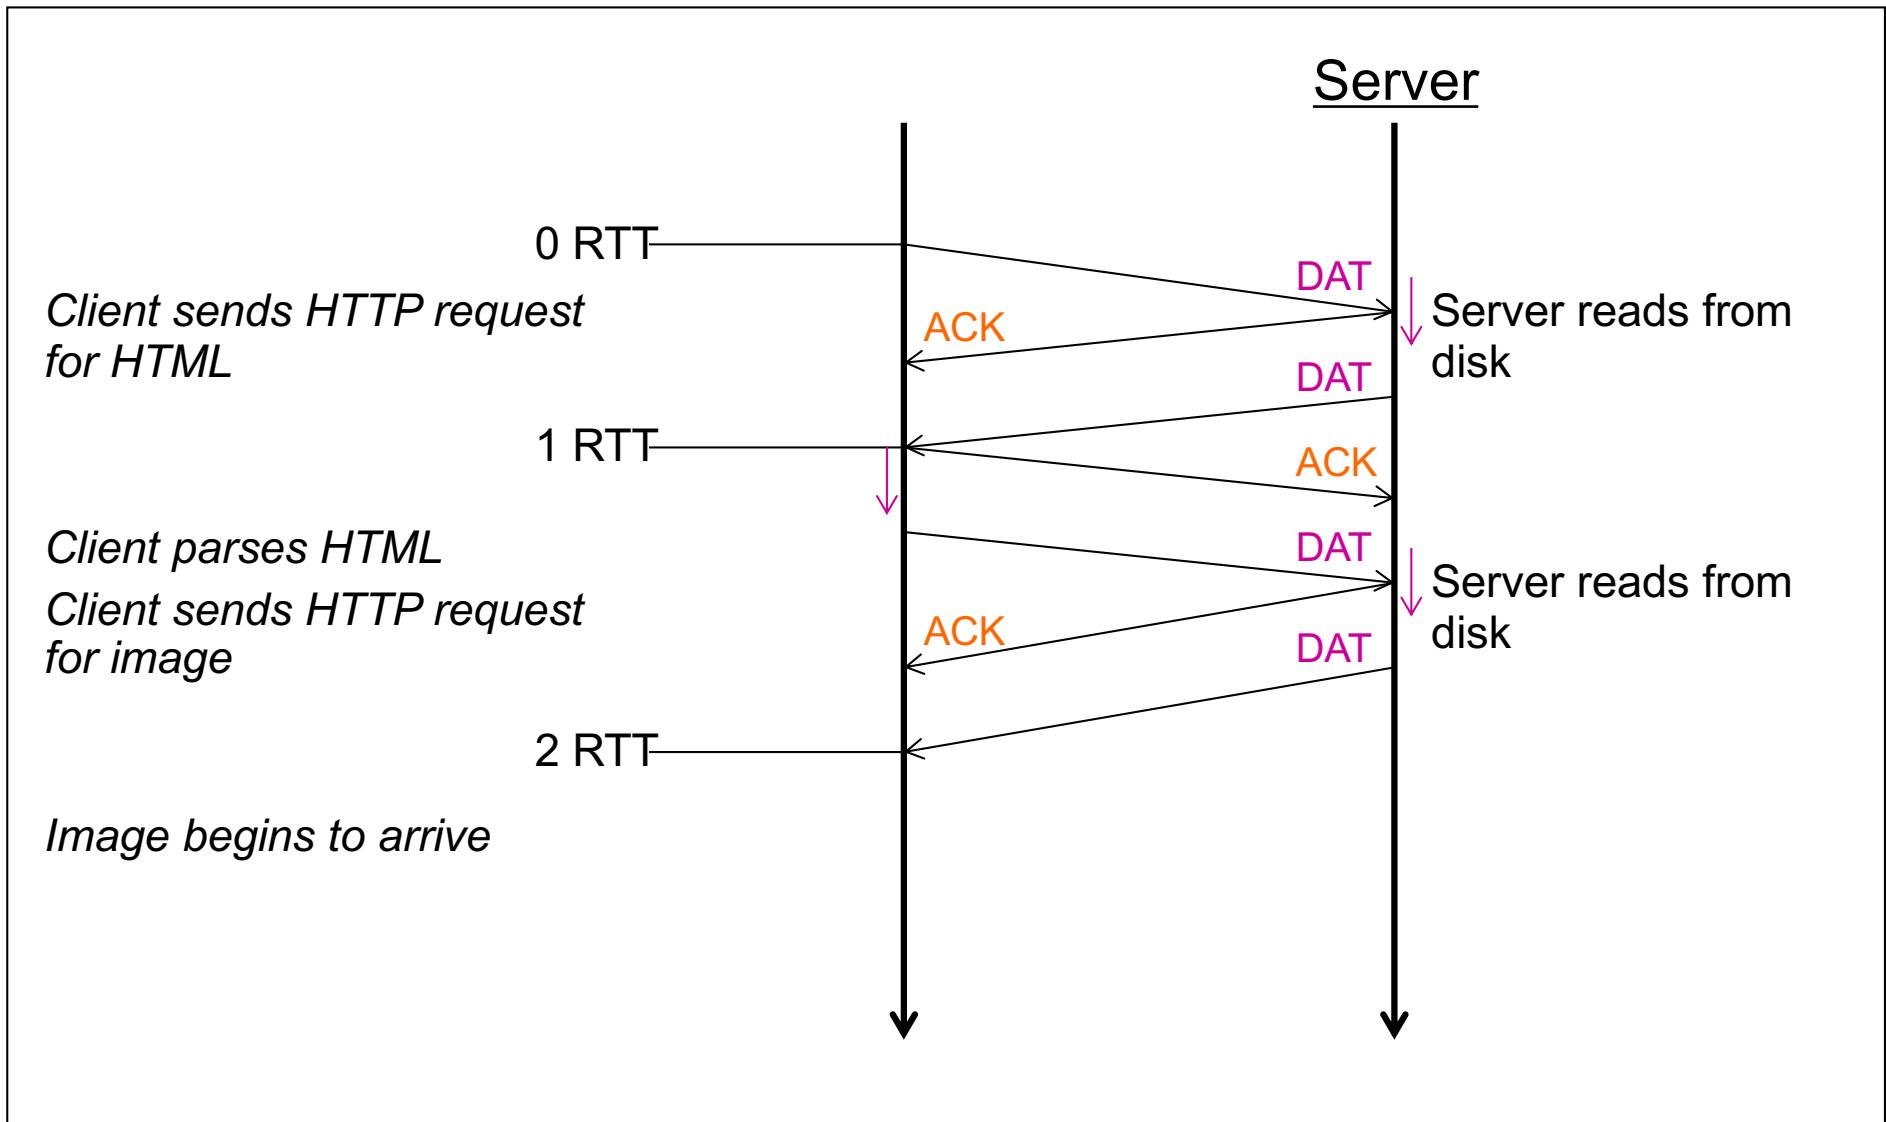
## Problems with simple model

- Multiple connection setups
  - Three-way handshake each time (TCP “synchronizing” stream)
- Lots of extra connections
  - Increases server state/processing
  - Server forced to keep connection state
- Later we will see also that
  - Short transfers are hard on stream protocol (TCP)
  - How much data should it send at once?
  - Congestion avoidance: Takes a while to “ramp up” to high sending rate (TCP “slow start”)
  - Loss recovery is poor when not “ramped up”



# Persistent Connection Example

Source: Freedman



# Persistent HTTP

## Non-persistent HTTP issues:

- Requires 2 RTTs per object
- OS must allocate resources for each TCP connection
- But browsers often open parallel TCP connections to fetch referenced objects

## Persistent HTTP:

- Server leaves connection open after sending response
- Subsequent HTTP messages between same client/server are sent over connection



# Persistent HTTP

## Persistent without pipelining:

- Client issues new request only when previous response has been received
- One RTT for each object

## Persistent with pipelining:

- Default in HTTP/1.1 spec
- Client sends requests as soon as it encounters referenced object
- As little as one RTT for all the referenced objects
- Server must handle responses in same order as requests

- Persistent without pipelining most common:  
When does pipelining work best?
- Multiple parallel requests or pipelined requests?

Source: Freedman (partial)



# HTTP Caching

- Clients often cache documents
  - When should origin be checked for changes?
  - Every time? Every session? Date?
- HTTP includes caching information in headers
  - HTTP 0.9/1.0 used: “Expires: <date>”; “Pragma: no-cache”
  - HTTP/1.1 has “Cache-Control”
  - “No-Cache”, “Private”, “Max-age: <seconds>”
  - “E-tag: <opaque value>”
- If not expired, use cached copy
- If expired, use condition GET request to origin
  - “If-Modified-Since: <date>”, “If-None-Match: <etag>”
  - 304 (“Not Modified”) or 200 (“OK”) response

Source: Freedman



# HTTP Conditional Request

GET / HTTP/1.1

Host: sns.cs.princeton.edu

User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X  
10.5; en-US; rv:1.9.2.13)

Connection: Keep-Alive

If-Modified-Since: Tue, 1 Feb 2011 17:54:18 GMT

If-None-Match: "7a11f-10ed-3a75ae4a"

HTTP/1.1 304 Not Modified

Date: Wed, 02 Feb 2011 04:01:21  
GMT

Server: Apache/2.2.3 (CentOS)

ETag: "7a11f-10ed-3a75ae4a"

Accept-Ranges: bytes

Keep-Alive: timeout=15, max=100

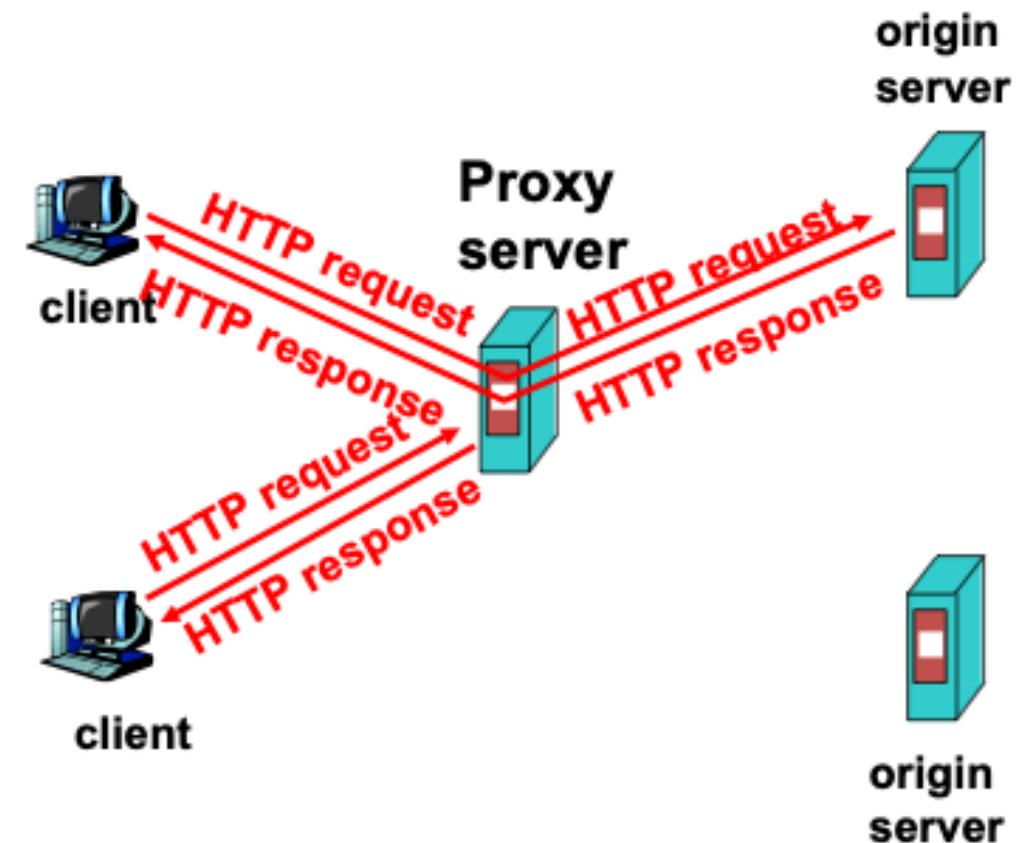
Connection: Keep-Alive



Source: Freedman

## Web Proxy Caches

- User configures browser: Web accesses via cache
- Browser sends all HTTP requests to cache
  - Object in cache cache returns object
  - Else: cache requests object from origin, then returns to client

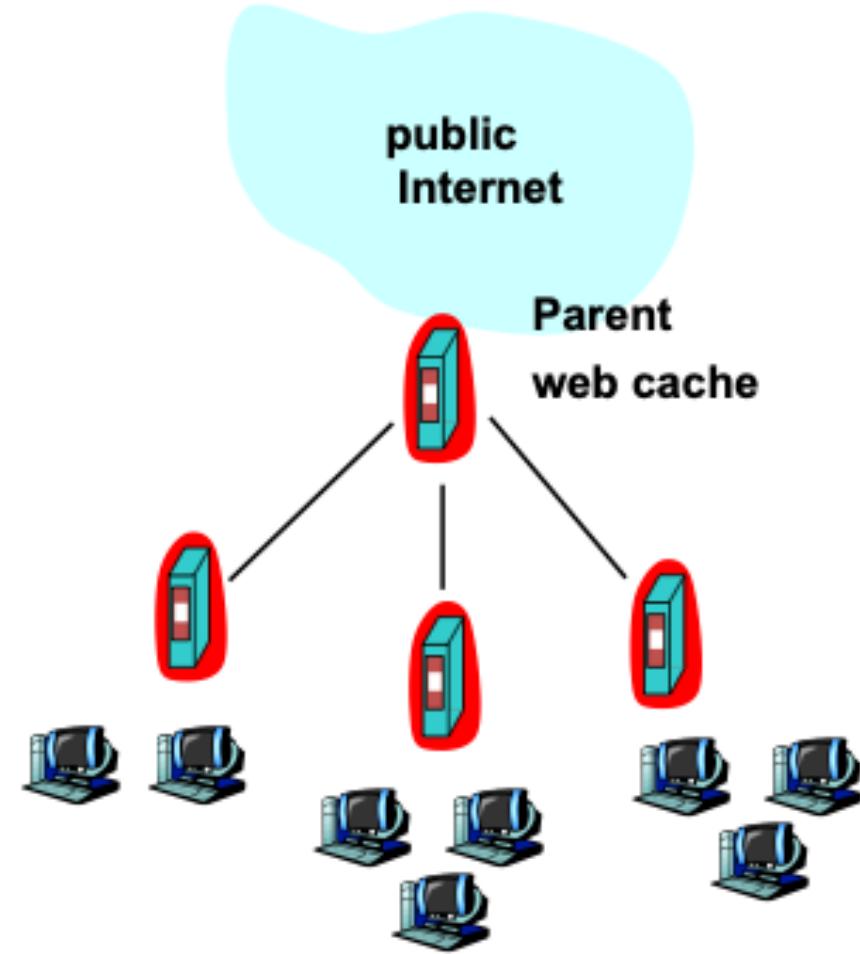


Source: Freedman



## When a single cache isn't enough

- What if the working set is > proxy disk?
  - Cooperation!
- A static hierarchy
  - Check local
  - If miss, check siblings
  - If miss, fetch through parent
- Internet Cache Protocol (ICP)
  - ICPv2 in RFC 2186 (& 2187)
  - UDP-based, short timeout



Source: Freedman



## Streaming multimedia: DASH

- *DASH: Dynamic, Adaptive Streaming over HTTP*
- *server:*
  - divides video file into multiple chunks
  - each chunk stored, encoded at different rates
- *manifest file:*
  - provides URLs for different chunks
- *client:*
  - periodically measures server-to-client bandwidth
  - consulting manifest, requests one chunk at a time
    - chooses maximum coding rate sustainable given current bandwidth
    - can choose different coding rates at different points in time (depending on available bandwidth at time)

Source: Kurose & Ross

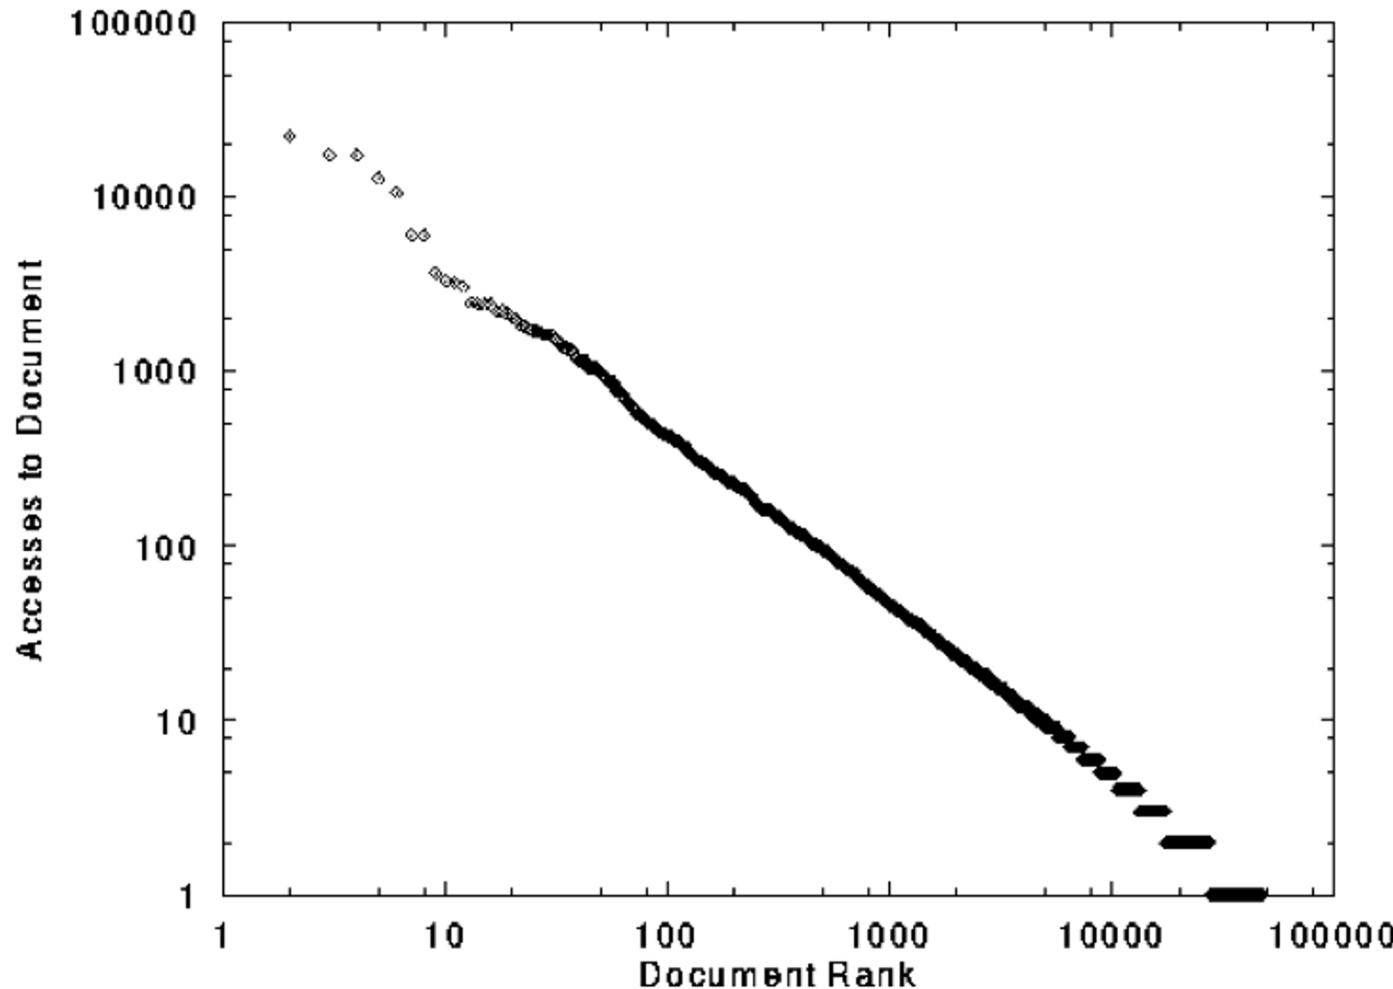


## Streaming multimedia: DASH

- *DASH: Dynamic, Adaptive Streaming over HTTP*
- “*intelligence*” at client: client determines
- *when* to request chunk (so that buffer starvation, or overflow does not occur)
- *what encoding rate* to request (higher quality when more bandwidth available)
- *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)



## Web traffic has cacheable workload



“Zipf” or “power-law” distribution

Source: Freedman

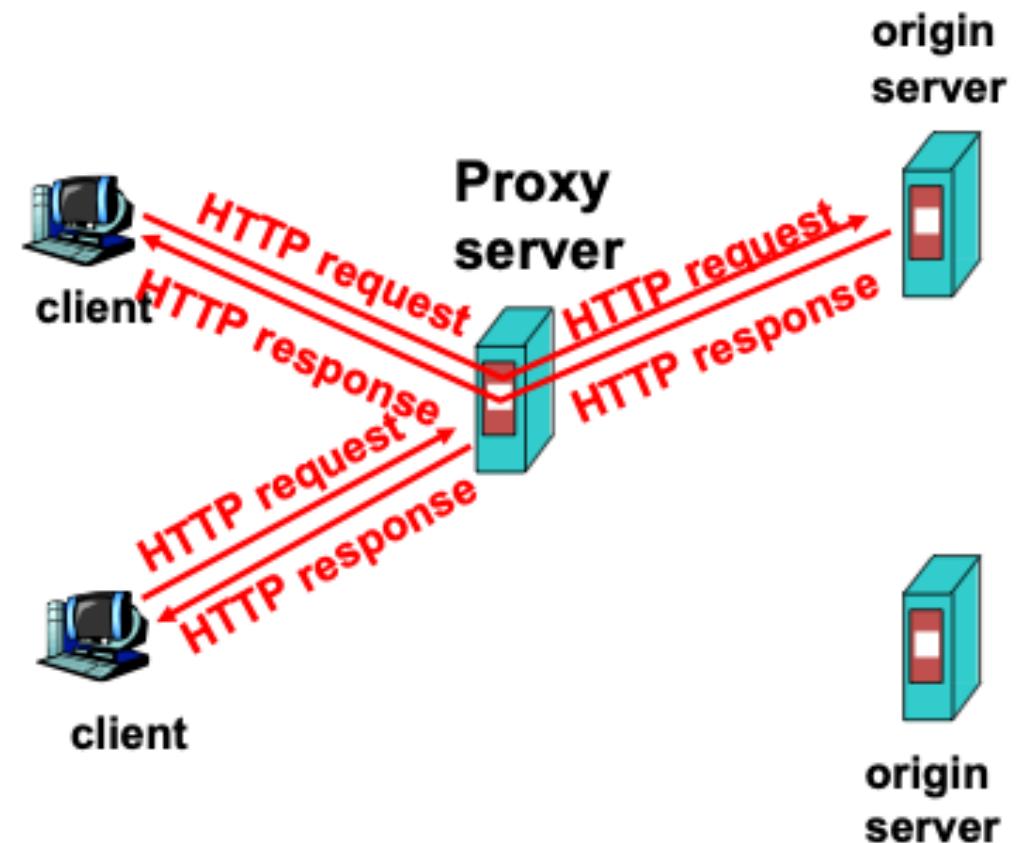


# Content Distribution Networks (CDNs)

- Content providers are CDN customers

## Content replication

- CDN company installs thousands of servers throughout Internet
  - In large datacenters
  - Or, close to users
- CDN replicates customers' content
- When provider updates content, CDN updates servers



Source: Freedman



# Content Distribution Networks & Server Selection

- Replicate content on many servers
- Challenges
  - How to replicate content
  - Where to replicate content
  - How to find replicated content
  - How to choose among known replicas
  - How to direct clients towards replica



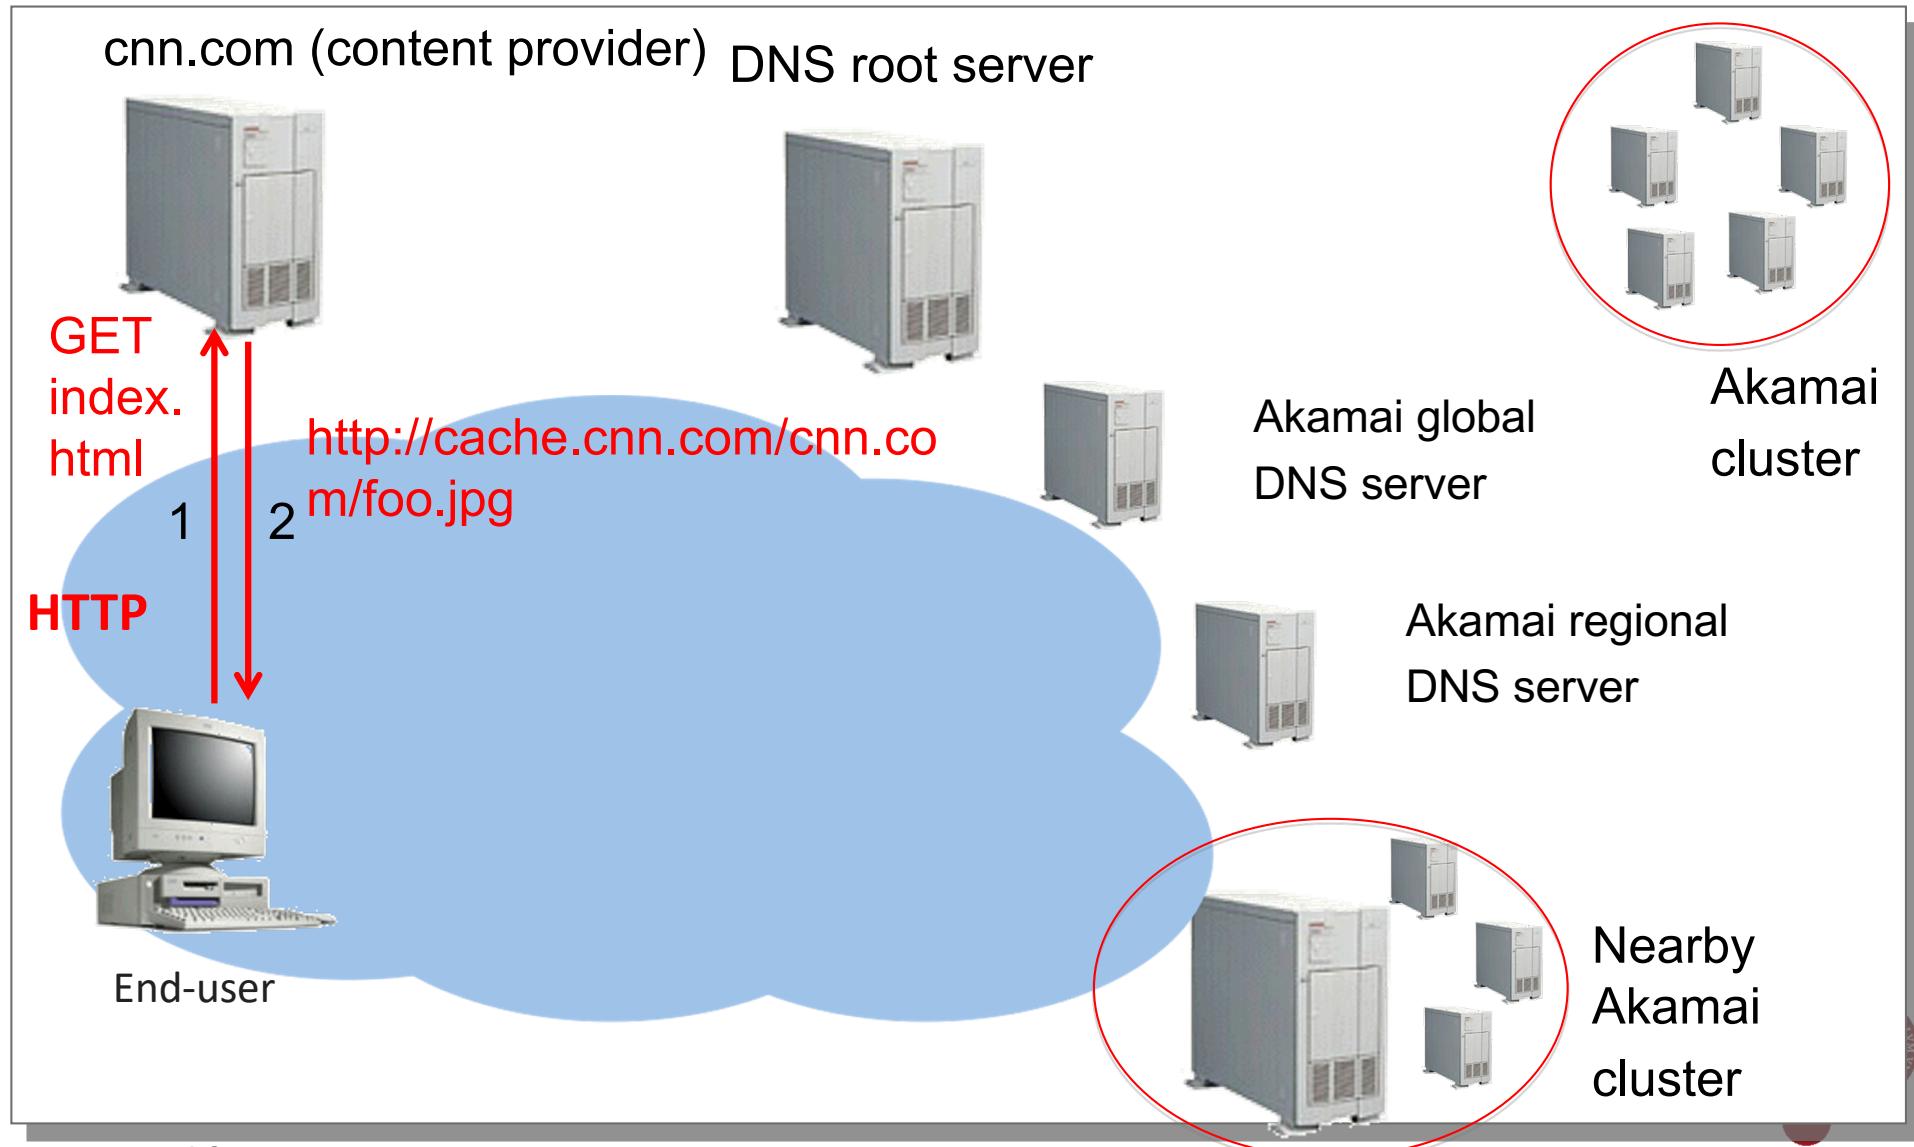
# Server Selection

- Which server?
- Lowest load: to balance load on servers
- Best performance: to improve client performance
- Based on Geography? RTT? Throughput? Load?
- Any alive node: to provide fault tolerance
  
- How to direct clients to a particular server?
- As part of routing: anycast, cluster load balancing
- As part of application: HTTP redirect
- As part of naming: DNS
  
- We will explain some of these techniques better later in the course!



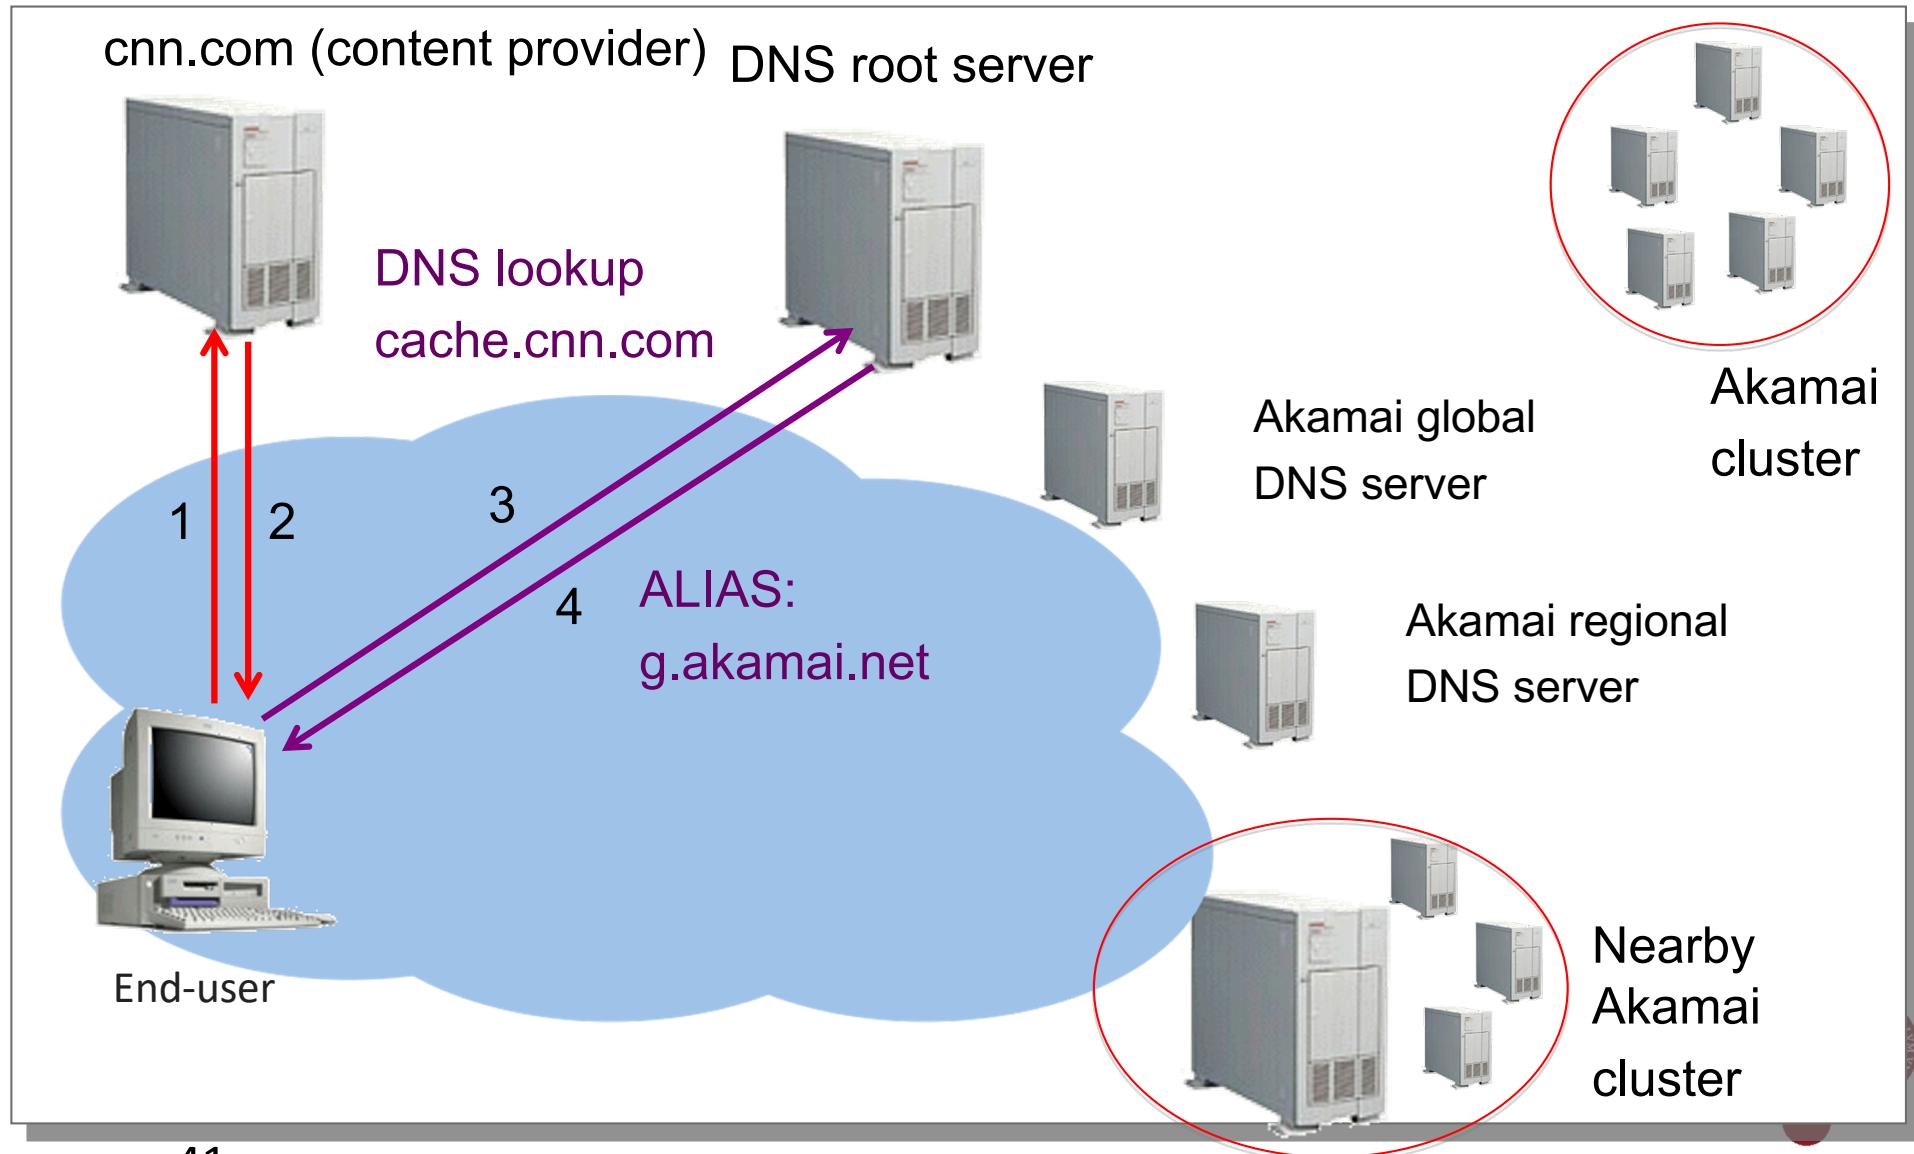
# How Akamai Works

Source: Freedman



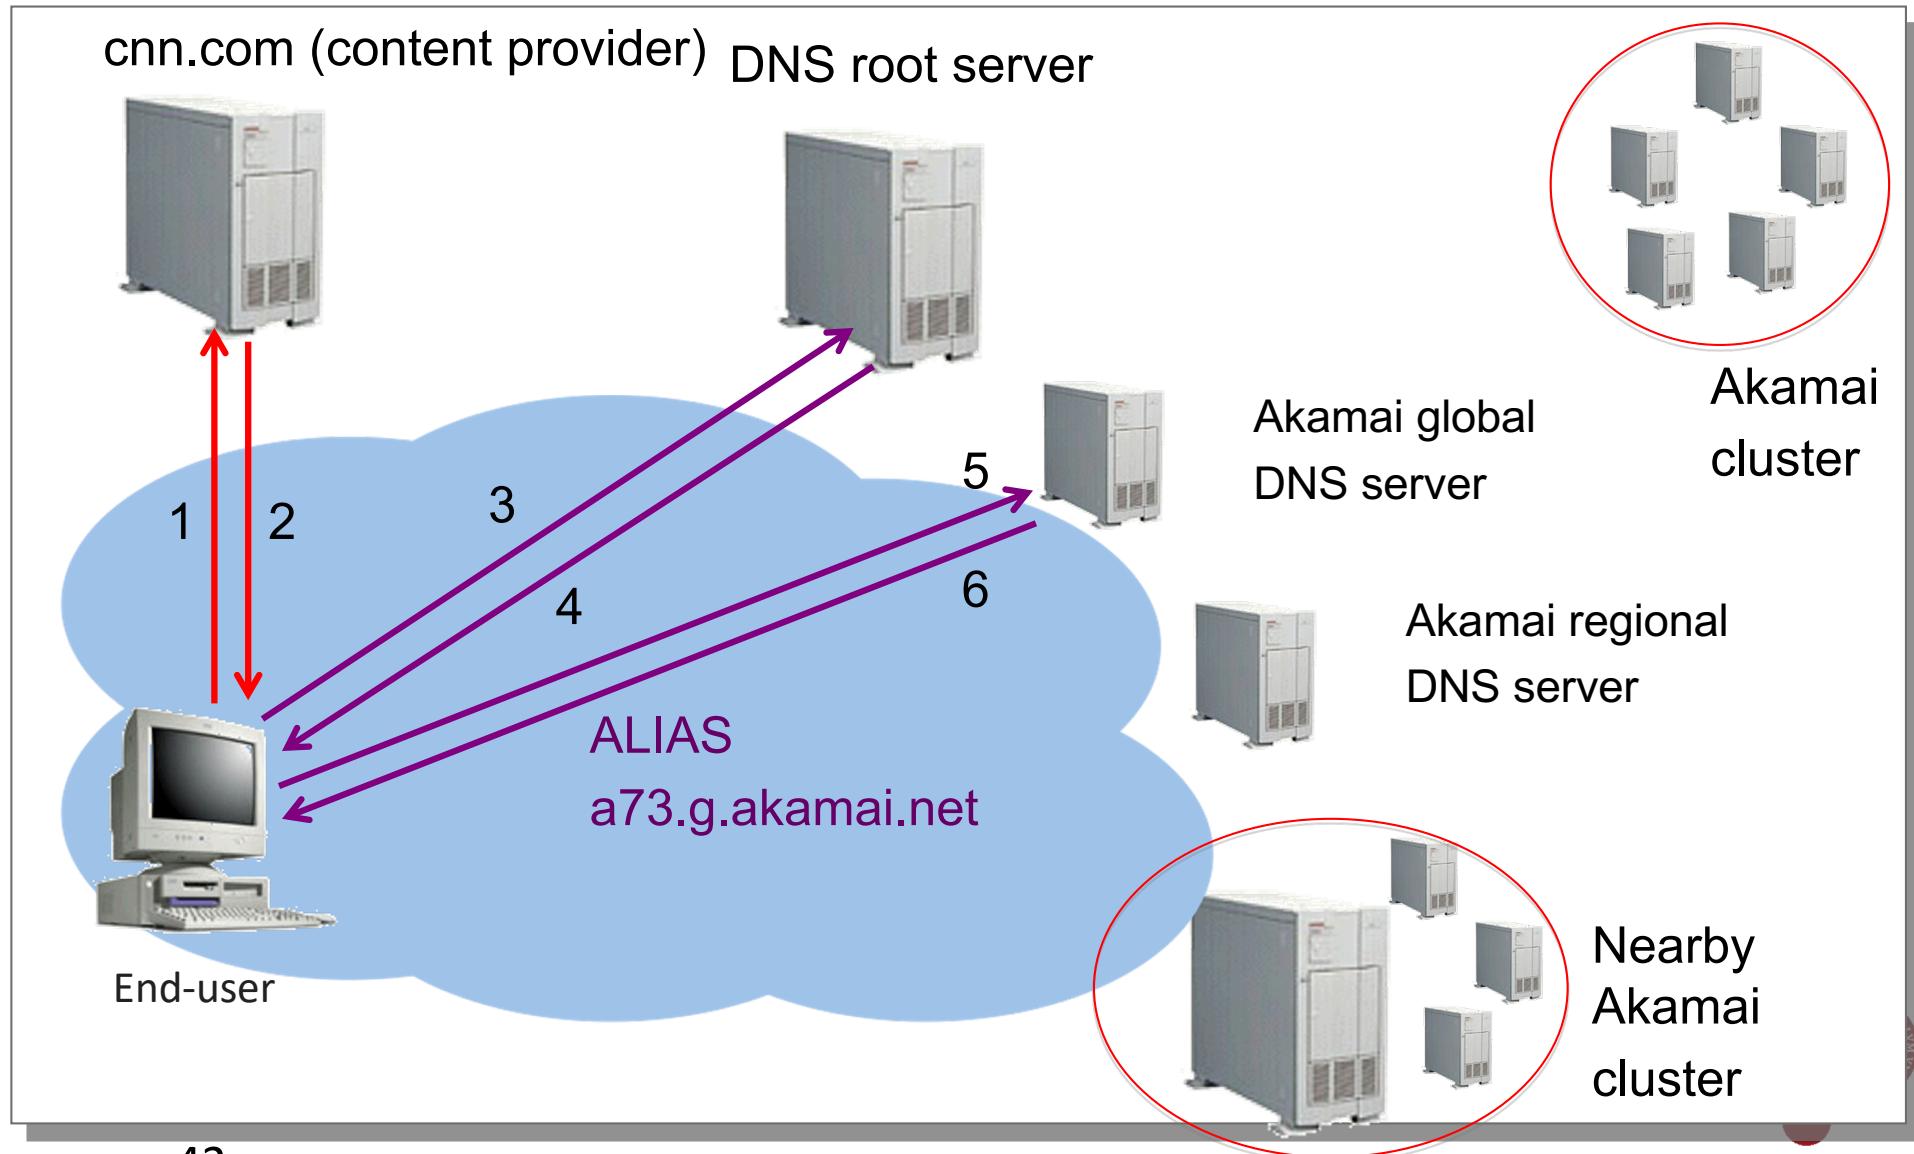
# How Akamai Works

Source: Freedman



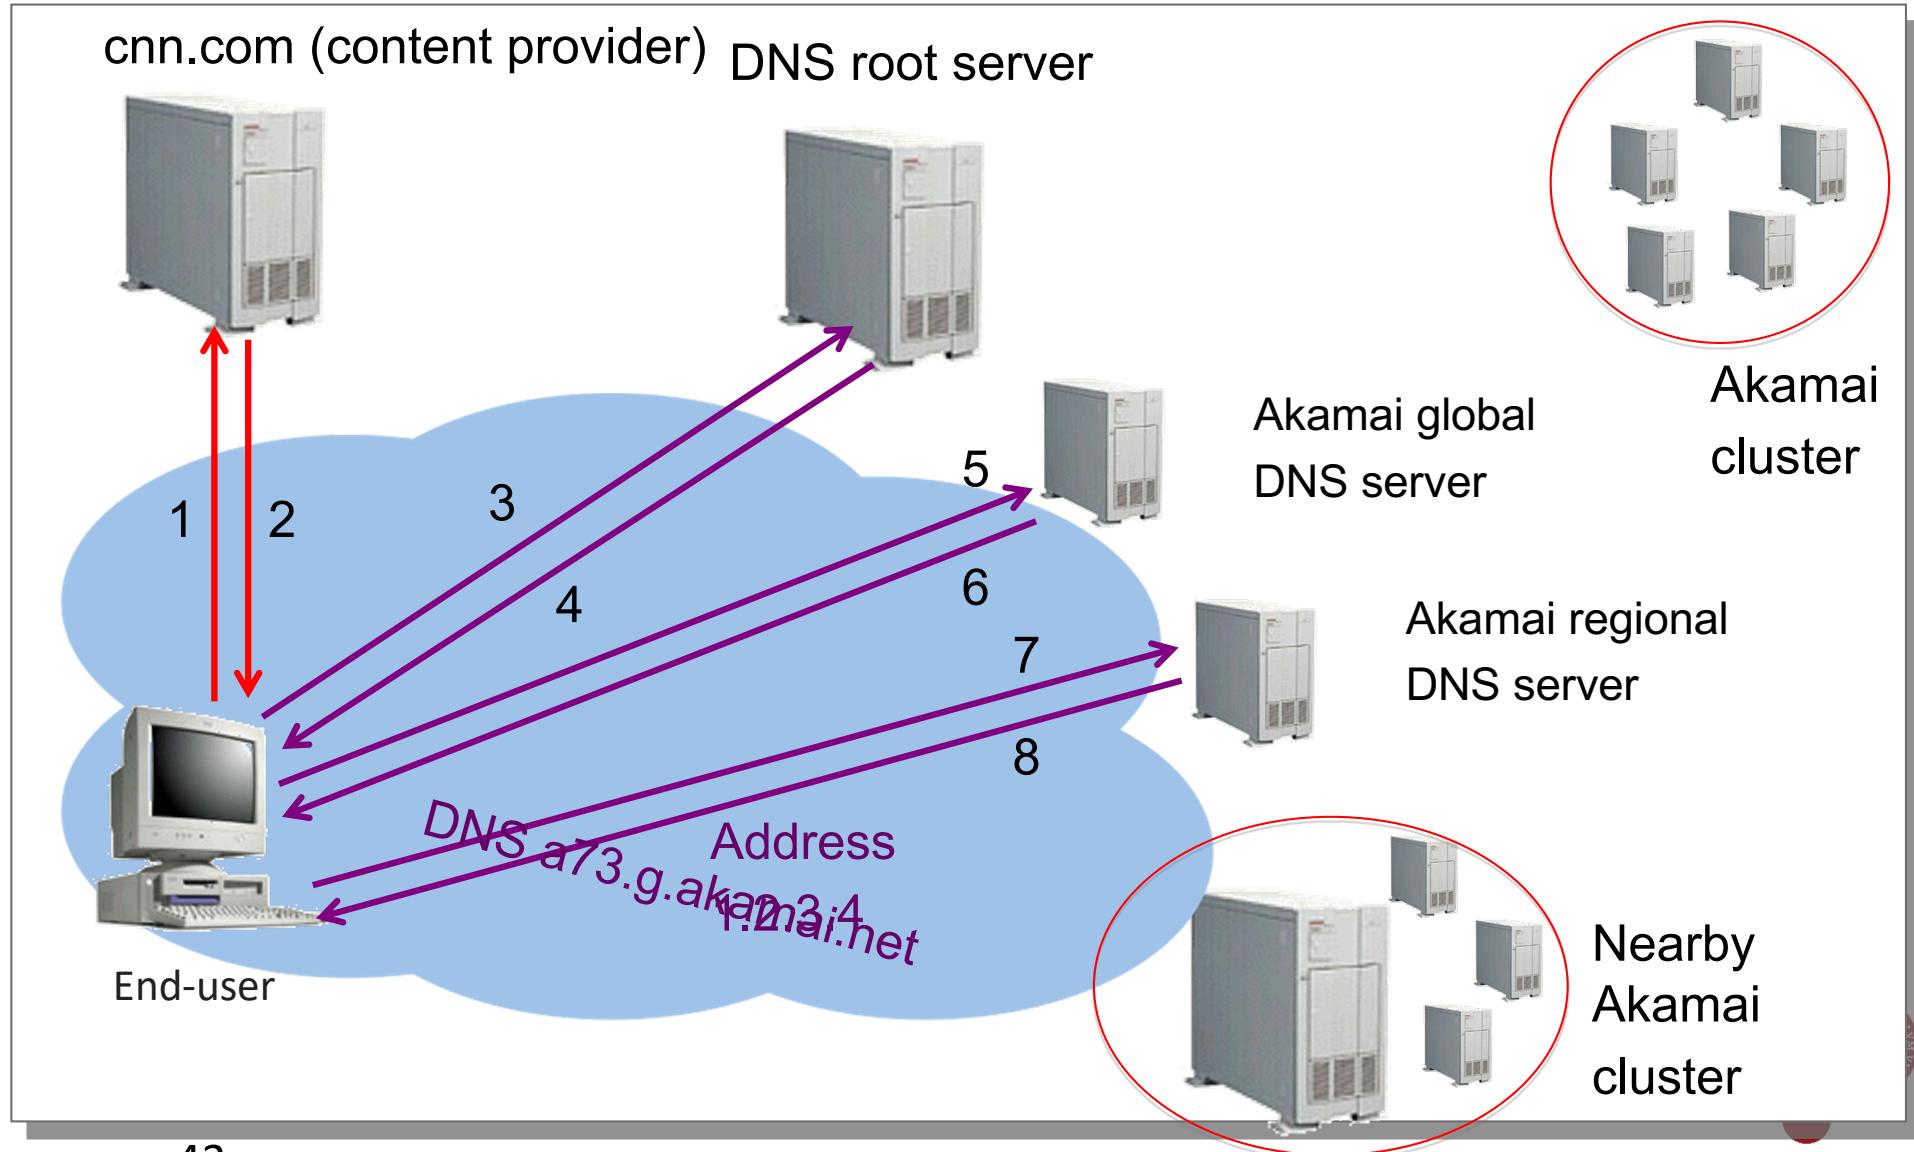
# How Akamai Works

Source: Freedman



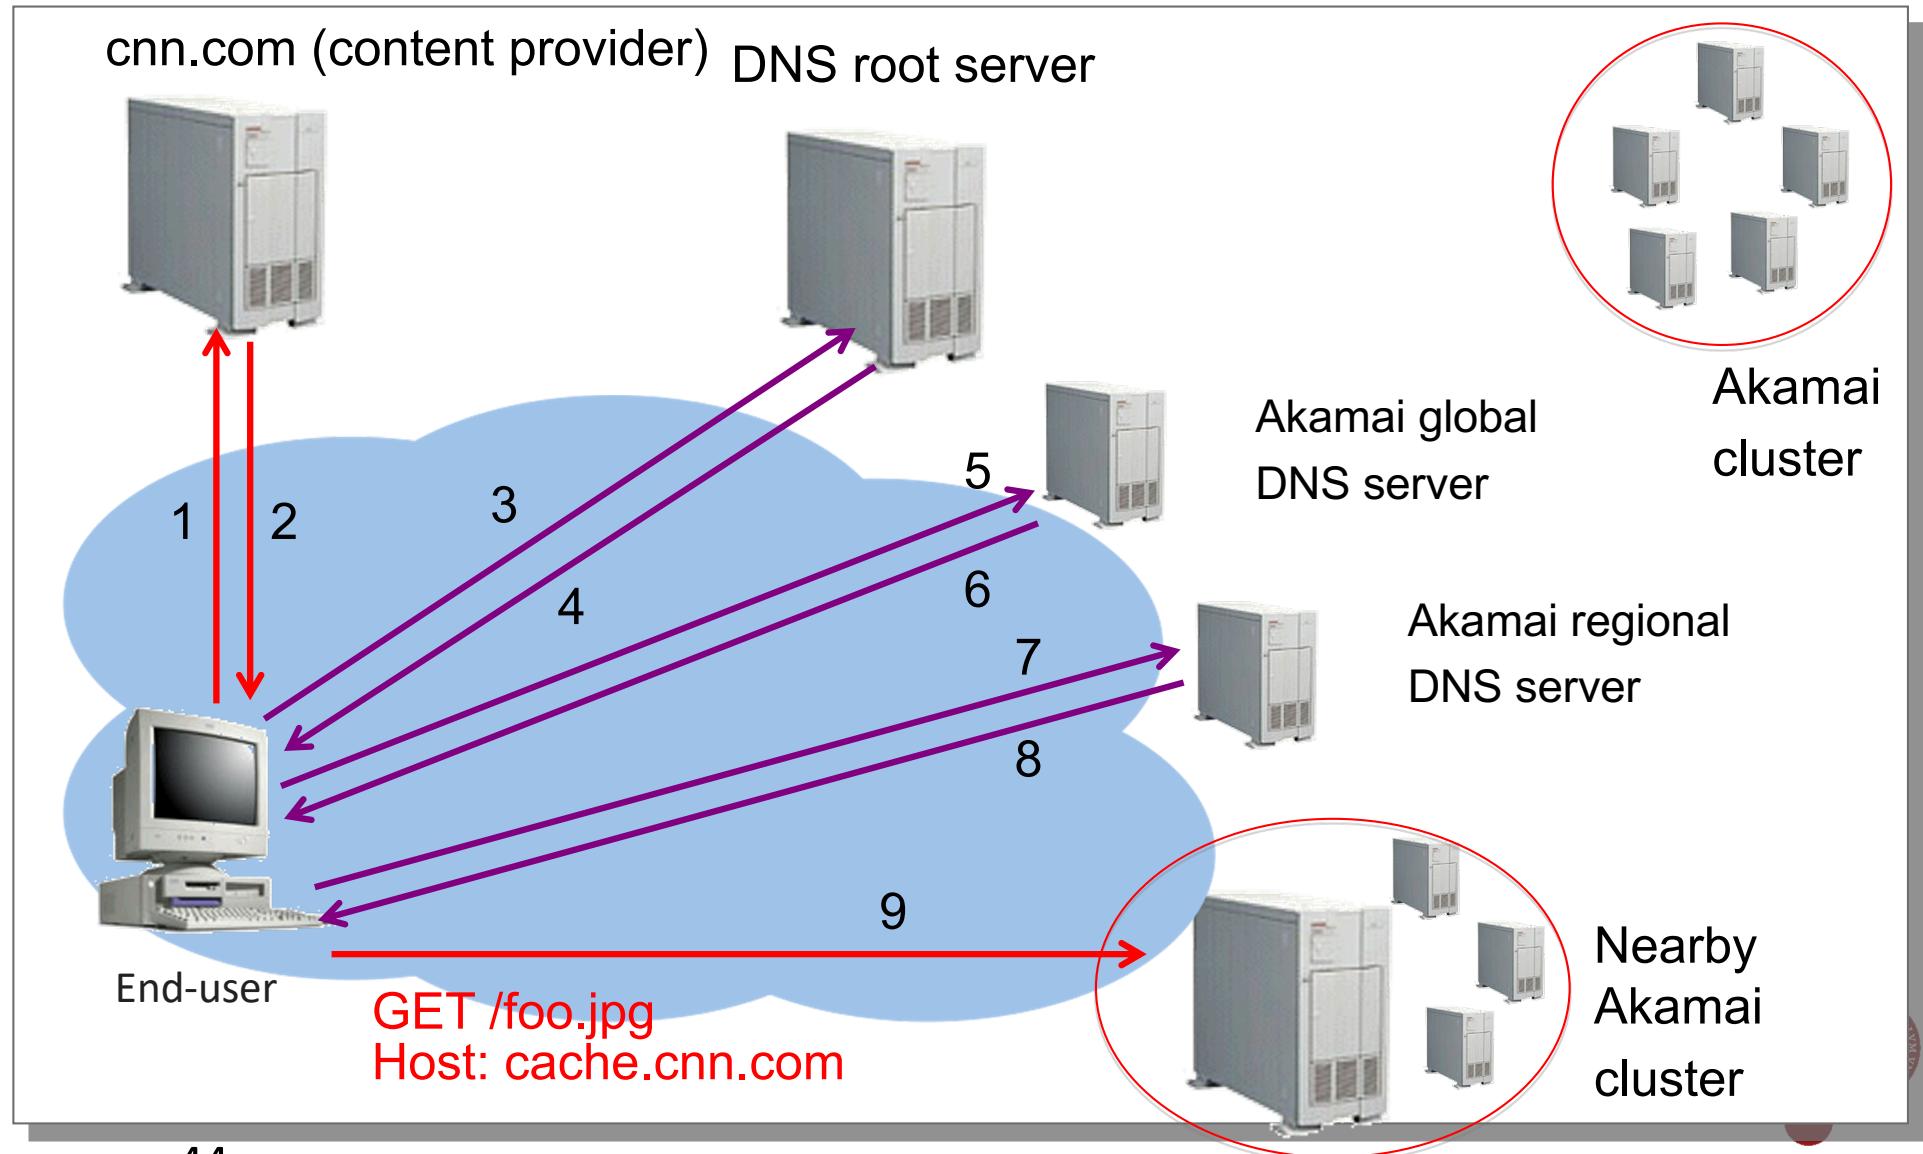
# How Akamai Works

Source: Freedman



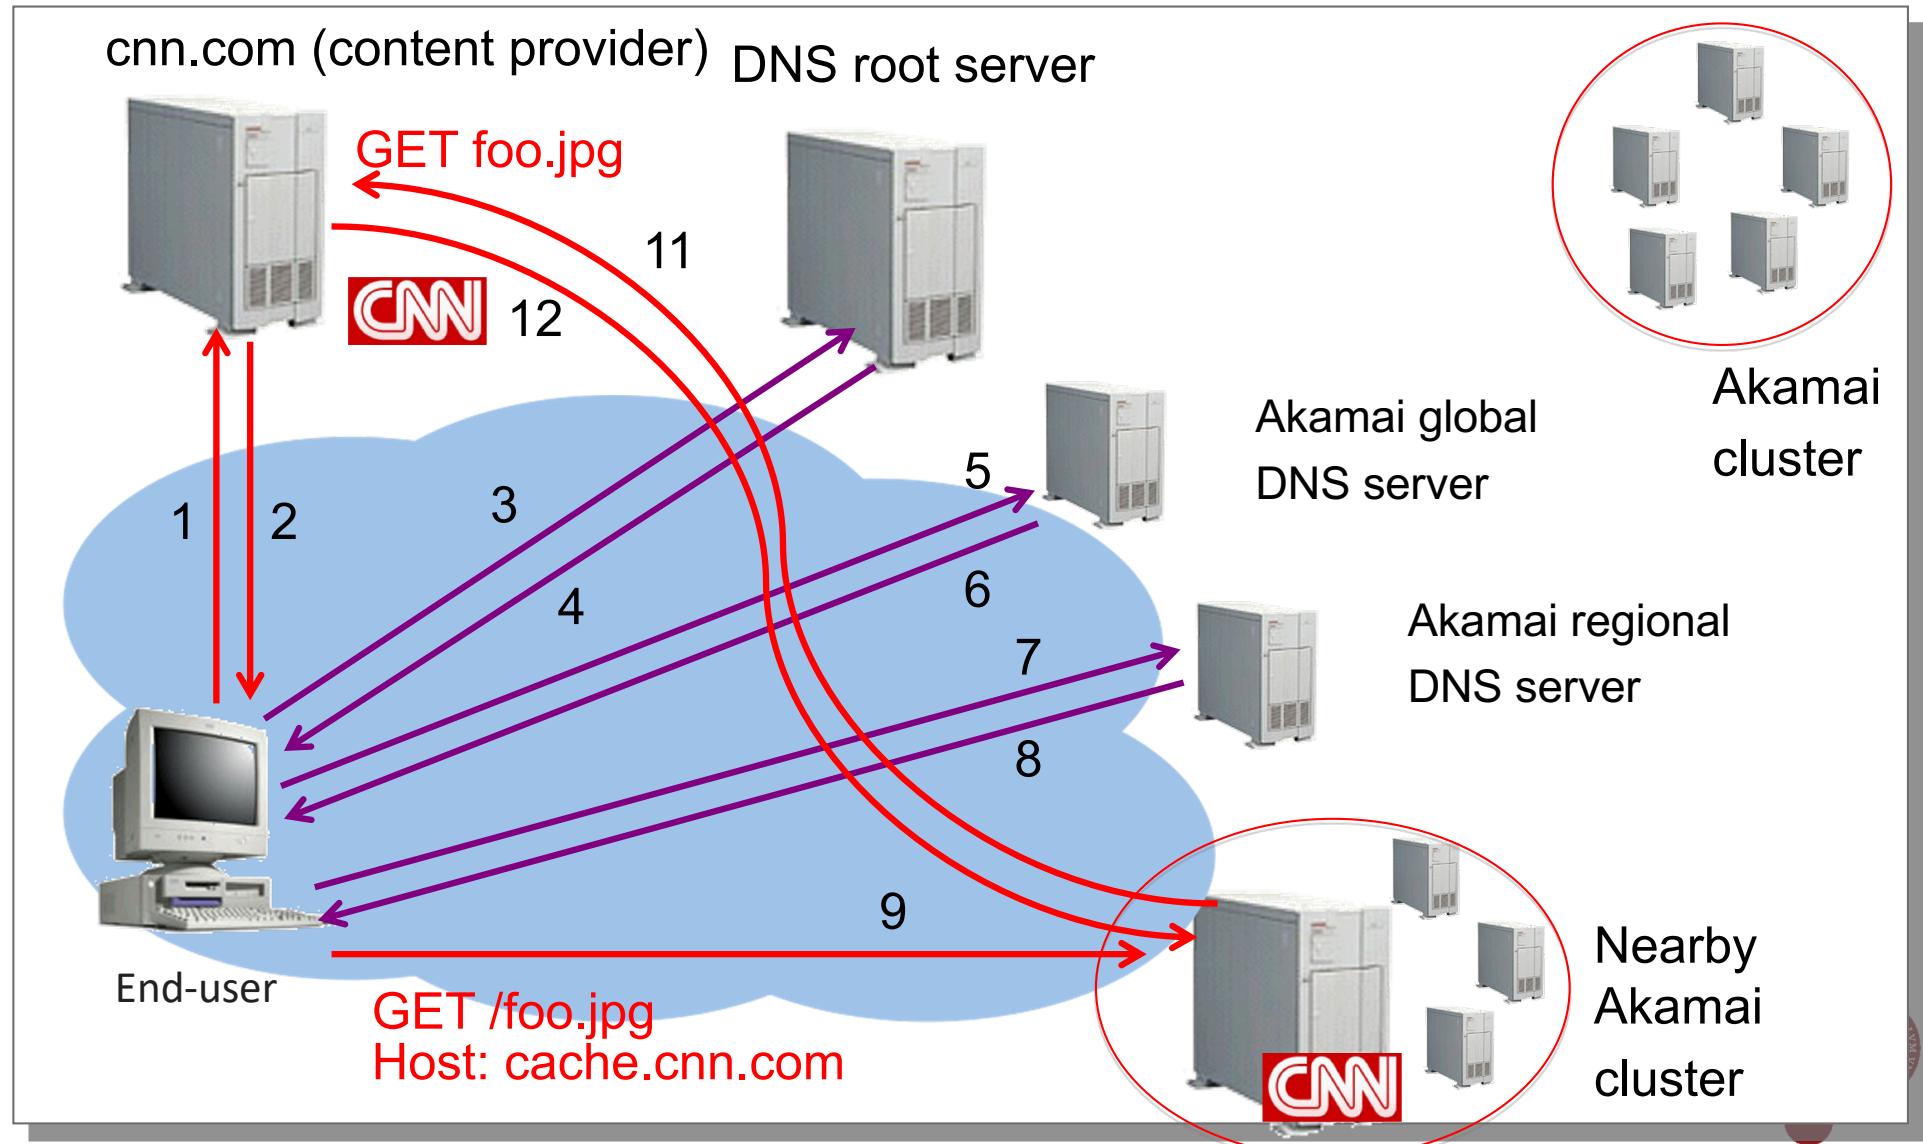
# How Akamai Works

Source: Freedman



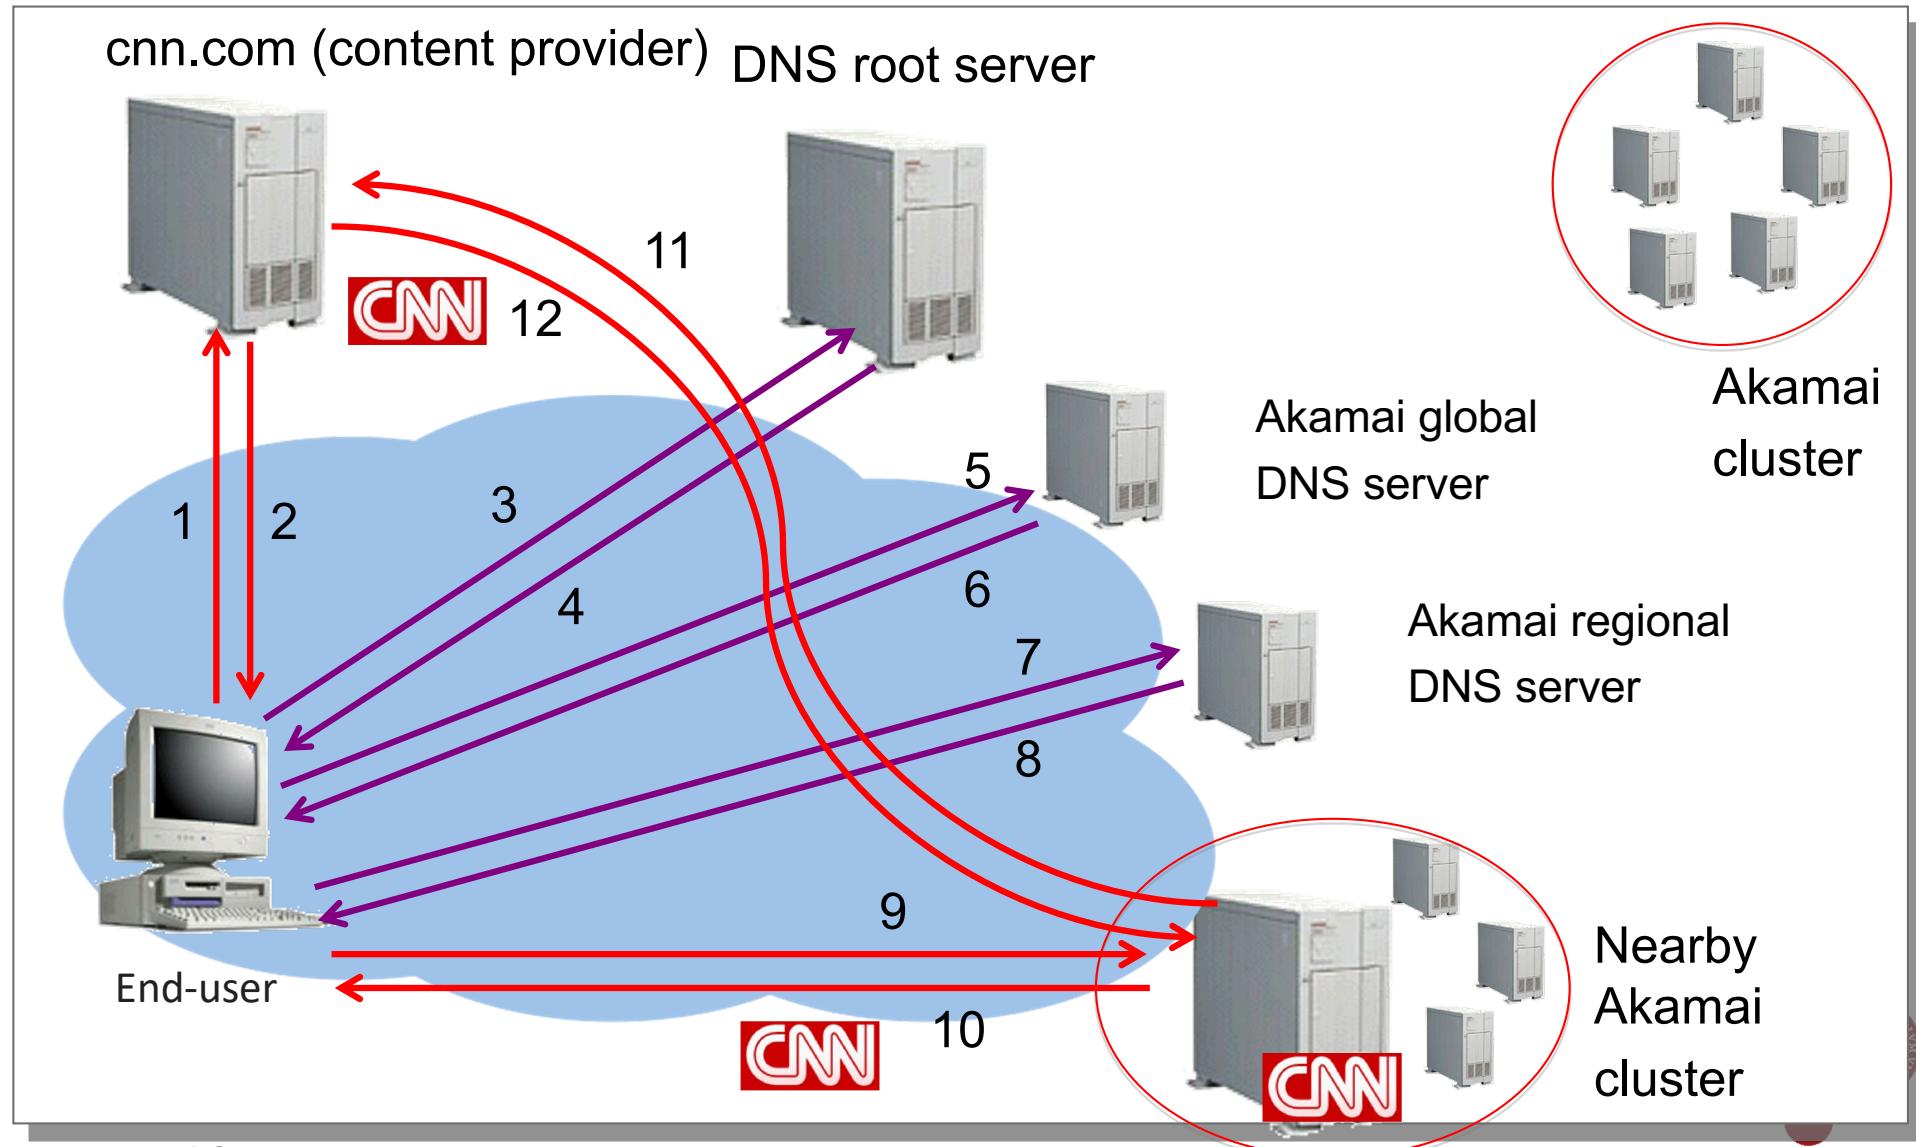
# How Akamai Works

Source: Freedman



# How Akamai Works

Source: Freedman



# Summary

- Network applications
- Email, Web → more in textbook, Chapter 2!
- Socket abstraction
- Communication between processes
- Client / Server, Peer-to-Peer
- HTTP concepts
- Web objects, request / response (pull)
- Persistent connections, web proxies
- Caching, content delivery
- Assignment A6 will be released for Wednesday

